

Historical perspective: Tensorflow version < 2

We will be using Tensorflow version 2 in this course.

- It has integrated the higher-level Keras API
- It uses "eager execution"

This notebook shows you

- The lower level non-Keras API
- Non-eager execution

The purpose

- It is interesting from an historical perspective
- Might give you an appreciation of Computation Graphs

Derived from Geron 11_deep_learning.ipynb

We will provide a quick introduction into programming with TensorFlow.

We revisit our old friend, MNIST digit classification and provide two solutions

- the first using "raw", low-level TensorFlow
- the second using the high-level Keras API

In [1]: USE_TF_VERSION=1

```
if USE_TF_VERSION < 2:  
    import tensorflow.compat.v1 as tf  
    tf.disable_v2_behavior()  
else:  
    import tensorflow as tf  
  
import numpy as np  
import os  
  
import pdb
```

WARNING:tensorflow:From /home/kjp/anaconda3/lib/python3.7/site-packages/tensorflow_core/python/compat/v2_compatible.py:65: disable_resource_variables (from tensorflow.python.ops.variable_scope) is deprecated and will be removed in a future version.

Instructions for updating:

non-resource variables are not supported in the long term

In [2]: print("Tensorflow version: ", tf.__version__)

Tensorflow version: 2.0.0

Raw TensorFlow

In []:

TensorFlow.layers

We will build an MNIST classifier using TensorFlow.layers

Get the MNIST dataset

- data presplit into training and test sets
 - flatten the images from 2 dimensional to 1 dimensional (makes it easier to feed into first layer)
 - create validation set from part of training
- "normalize" the inputs: change pixel range from [0,255] to [0,1]

```
In [3]: (X_train, y_train), (X_test, y_test) = tf.keras.datasets.mnist.load_data()

# Determine
# - the dimensions of the input by examining the first training example
# - the dimensions of the output (number of classes) by examining the targets
input_size = np.prod(X_train[0].shape)
output_size = np.unique(y_train).shape[0]

# input image dimensions
img_rows, img_cols = X_train[0].shape[0:2]

valid_size = X_train.shape[0] // 10

# Flatten the data to one dimension and normalize to range [0,1]
X_train = X_train.astype(np.float32).reshape(-1, input_size) / 255.0
X_test = X_test.astype(np.float32).reshape(-1, input_size) / 255.0
y_train = y_train.astype(np.int32)
y_test = y_test.astype(np.int32)
X_valid, X_train = X_train[:valid_size], X_train[valid_size:]
y_valid, y_train = y_train[:valid_size], y_train[valid_size:]
```

```
In [4]: X_train.shape
```

```
Out[4]: (54000, 784)
```

```
In [5]: (n_hidden_1, n_hidden_2) = (100, 30)
```

```
In [6]: # Placeholders for input X, target y
# The first dimension (None) is for the batch size

X = tf.placeholder(tf.float32, shape=(None, input_size), name="X")
y = tf.placeholder(tf.int32, shape=(None), name="y")
```

Create function to return mini-batches

```
In [7]: def next_batch(X, y, batch_size, shuffle=True):
```

```
    """
```

```
        Generator to return batches from X and y
```

Parameters

```
    -----
```

X: ndarray

y: ndarray. The first dimension of X and y must be the same

batch_size: Int. The size of the slice (of X and y) to return in each batch

shuffle: Boolean. Sample X, y in random order if True

Yields

```
    -----
```

X_batch, y_batch: a 2-tuple of ndarrays,

- where X_batch is a slice (of size at most batch_size) of X

- where y_batch is a slice of y (same first dimension as X_batch)

If first dimension of X is not evenly divisible by batch size, the final batch will

be of size smaller than batch_size

```
    """
```

```
# Randomize the indices
```

```
if shuffle:
```

```
    idx = np.random.permutation(len(X))
```

```
else:
```

```
    idx = np.arange( len(X) )
```

```
# Return a batch of size (at most) batch_size,
```

```
# starting at idx[next_start]
```

```
next_start = 0
```

```
n_batches = len(X) // batch_size
```

```
while next_start < len(X):
```

```
# Get a batch of indices from idx, starting at idx[next_start] and ending at idx[next_end]
    next_end = min(next_start + batch_size, len(X))
    X_batch, y_batch = X[ idx[next_start:next_end] ], y[ idx[next_start:next_end] ]

# Advance next_start to start of next batch
next_start = next_start + batch_size

# Return a batch
yield X_batch, y_batch
```

Build the computation graph

```
In [8]: (n_hidden_1, n_hidden_2) = (100, 30)
```

```
In [9]: # to make this notebook's output stable across runs
def reset_graph(seed=42):
    tf.reset_default_graph()
    tf.set_random_seed(seed)
    np.random.seed(seed)
```

In [10]: `reset_graph()`

```
# Placeholders for input X, target y
# The first dimension (None) is for the batch size
X = tf.placeholder(tf.float32, shape=(None, input_size), name="X")
y = tf.placeholder(tf.int32, shape=(None), name="y")

with tf.name_scope("dnn"):
    hidden1 = tf.layers.dense(X, n_hidden_1, activation="relu", name="hidden1")
    hidden2 = tf.layers.dense(hidden1, n_hidden_2, activation="relu", name="hidden2")
    logits = tf.layers.dense(hidden2, output_size, name="outputs_")
```

WARNING:tensorflow:From <ipython-input-10-7a8e8474ba41>:10: dense (from tensorflow.python.layers.core) is deprecated and will be removed in a future version.

Instructions for updating:

Use keras.layers.Dense instead.

WARNING:tensorflow:From /home/kjp/anaconda3/lib/python3.7/site-packages/tensorflow_core/python/layers/core.py:187: Layer.apply (from tensorflow.python.keras.engine.base_layer) is deprecated and will be removed in a future version.

Instructions for updating:

Please use `layer.__call__` method instead.

Create a loss node

- Use cross entropy as loss
 - we are comparing the probability vector computed by the graph (logits) with the target probability vector (y)

Ordinarily we would need to

- convert the scores (logits) vector to a probability vector by a softmax activation on the "outputs" layer
- convert the target to a one-hot vector (length equal to number of target classes, which is also length of probability vector)
- compare the two vectors with cross_entropy

TensorFlow provides a very convenient method

`sparse_softmax_cross_entropy_with_logits` that does all the work for us !

- applies softmax to the scores (logits)
- converts integer targets (in range [0, number of classes]) into one-hot vectors (with length equal to number of classes)
- ~~does the one-hot encoding~~

```
In [11]: with tf.name_scope("loss"):  
    # xentropy is a tensor whose first dimension is the batch size  
    xentropy = tf.nn.sparse_softmax_cross_entropy_with_logits(labels=y, logits=logits)  
  
    # Find the loss across the examples in the batch by summing individual example losses  
    loss = tf.reduce_mean(xentropy, name="loss")
```

Create a node to compute accuracy

- for each example, compares the element in the logit vector with the highest score (i.e., index of our prediction) to the target
- sums up the number of examples with matching max logit and target

```
In [12]: with tf.name_scope("eval"):  
    correct = tf.nn.in_top_k(logits, y, 1)  
    accuracy = tf.reduce_mean(tf.cast(correct, tf.float32))
```

Create the training operations

- Training operation is an optimizer step that minimizes the loss

```
In [13]: learning_rate = 0.01  
  
with tf.name_scope("train"):  
    optimizer = tf.train.GradientDescentOptimizer(learning_rate)  
    training_op = optimizer.minimize(loss)
```

Create an initialization node to initialize global variables (i.e., the weights that the optimizer will solve for)

```
In [14]: init = tf.global_variables_initializer()  
saver = tf.train.Saver()
```

Run the training loop

- Run for multiple "epochs"; an epoch is an entire pass through the training data set
- For each epoch, divide the training set into mini-batches
 - For each mini-batch
 - run the "training operation" (i.e, the optimizer)
 - every few epochs
 - compute the accuracy (by evaluating the graph node that computes accuracy) on the training and validation set

In general, we usually continue training

- as long as the validation loss continues to decrease across epochs

In [15]:

```
n_epochs = 20  
batch_size = 50  
  
modelName = "mnist_first"  
  
save_path = os.path.join(".", modelName + ".ckpt")
```

