

# Policy-based methods/Policy gradient methods

## Value-based methods

- assign a value to a state (value function) OR action given a state (action value function)
- policy is *derived* from these values
  - chose the action leading to the greatest return
  - $\text{\argmax}$  \act to implement the policy

## Policy-based methods

- by contrast, construct the policy  $\pi$  directly
- as a parameterized (by parameters  $\Theta$ ) function

$$\pi_\theta(\text{\actseq} | \text{\stateseq}) = \text{\prc} \text{\actseq}_= \text{\act} \text{\stateseq}_= \text{\state}$$

i.e., the policy is a probability distribution of actions \act, conditional on the state \state)

Policy-based methods are *necessary* in those cases in which Value-based methods are not possible:

- Continuous (versus discrete) action
  - the  $\text{\argmax}$  \act that implements policy in Value based methods is not possible
- Stochastic policy necessary
  - games against an adversary: when an adversary can take advantage of Agent predictability

Policy-based methods are *desirable/preferable* when Value-based methods are impractical

- Large number of possible actions
- High dimensional state spaces
  - state is characterized by a (long) vector of characteristics

In both these cases:

- tables are impractical representations of Value function or Action value function

Scenario	Policy-Based Required	Policy-Based Desirable
Continuous action spaces	Yes	Yes
Stochastic strategies needed	Yes	Yes
Aliased or partially observable states	Yes	Yes
High-dimensional spaces	Sometimes	Yes
Discrete/simple environments	No	Sometimes

Here is a brief comparison of Value-based and Policy-based methods.

Aspect	Value-Based Methods	Policy-Based Methods
Output	State/action value functions	Directly parameterized policy
Policy Representation	Implicit (via greedy/exploratory actions)	Explicit (probability/distribution mapping)
Learning Objective	Value prediction loss minimization	Expected return maximization (gradient ascent)
Typical Example Algorithms	DQN, Q-learning, SARSA	REINFORCE, PPO, vanilla policy gradient
Action Space	Discrete (practical)	Handles continuous and discrete
Stochastic Policies	Limited	Natural/efficient
Exploration Strategies	Decoupled from policy (e.g. epsilon-greedy)	Inherent (stochastic policy outputs)

# Policy Gradient methods

The predominant class of Policy based methods are those based on the *Policy Gradient* method.

Policy Gradient methods create a sequence of improving policies

$$\pi_0, \dots, \pi_p, \dots$$

by creating a sequence of improved parameter estimates

$$\theta_0, \dots, \theta_p, \dots$$

using Gradient Ascent on some objective function  $J(\theta)$  to improve  $\theta_p$

$$\theta_{p+1} = \theta_p + \alpha * \nabla_{\theta} J(\theta_p)$$

- gradient of an Objective Function  $J(\theta)$
- with respect to parameters  $\Theta$

Since we are trying to maximize objective function  $J(\theta)$  rather than minimize a loss objective

- we use Gradient Ascent rather than Gradient Descent
- hence we add the gradient rather than subtract it, in the update

RL Book Chapt 12 (<http://incompleteideas.net/book/RLbook2020.pdf#page=343>)

## Aside

There are a few Policy base methods that *don't* use Policy Gradient

- in the module on Value based methods, we learned about Policy Iteration
- Policy iteration alternates
  - Policy Evaluation: improving the estimate of a Value function
  - Policy Improvement: improving the policy
    - use  $\text{\argmax}$  \act to implement the current policy

Thus, Policy Iteration is both Value based and Policy based

- but does not evolve policy via gradients

# Stochastic policy and environment

With Value based methods

- the Environment can be stochastic
- but the Policy is usually deterministic
  - $\text{\argmax}$  \act to implement the policy

With Policy Gradient methods

- the policy can be stochastic (action is a probability distribution)  
$$\pi(\text{\act} | \text{\state}; \theta) = \text{\pr}\text{\actseq} = \text{\act} | \text{\stateseq} = \text{\state}, \theta = \theta$$
- for example: Non-greedy policy that trades off Exploitation vs Exploration

The environment can *also* be stochastic \$\$

$\text{\textbackslash transp}(\{ \text{\textbackslash state}', \text{\textbackslash rew} | \text{\textbackslash state}, \text{\textbackslash act } \})$

$\text{\textbackslash transp}(\{ \text{\textbackslash stateseq}\{tt+1}, \text{\textbackslash rewseq}\{tt+1} | \text{\textbackslash state} = \text{\textbackslash stateseq}\{tt}, \text{\textbackslash act} = \text{\textbackslash actseq}\{tt } \})$  \$\$

- the response ( $\text{\textbackslash state}'$ ,  $\text{\textbackslash rew}$ ) by the environment is not deterministic

This poses a challenge to Value-based methods

- a single observation of ( $\text{\textbackslash state}'$ ,  $\text{\textbackslash rew}$ ) is a *high variance* estimate of  
 $\text{\textcolor{red}{\text{\textbackslash transp}(\text{\textbackslash state}', \text{\textbackslash rew}| \text{\textbackslash state}, \text{\textbackslash act})}}$

# Objective function

The performance measure  $J(\theta)$  that we seek to maximize is the

- *expected value* (across each possible episode  $\tau$ ) of
- the return  $G_{0,\tau}$  from initial state  $\text{\textbackslash stateseq}_0$  of the episode

$$J(\theta) = \text{\textbackslash Exp} \tau \sim \text{\textbackslash pr} \theta(G_{0,\tau})$$

- where  $\text{\textbackslash pr} \theta$  is the probability distribution of episodes
  - is a function of the policy parameters  $\theta$

Recall: the return from state  $\text{\textbackslash stateseq\_tt}$  of the episode is

$$\begin{aligned} G_{,\tau} &= \sum_{k=0}^{\gamma^k * \text{\textbackslash rewseq}_{+k+1}} \\ &= \text{\textbackslash rewseq}_{+1} + \gamma * G_{+1,\tau} \end{aligned}$$

# Taking the gradient of the Objective

For Gradient Ascent/Descent

- We need to be able to compute

$$\nabla_{\theta} J(\theta_p)$$

the gradient of the Objective w.r.t the parameters

However: there is an issue in computing  $J(\theta)$ .

- Letting  $\text{pr}\tau; \theta$  denote the probability of episode  $\tau$  occurring
- the expectation can be re-written as a probability-weighted sum

$$\text{\textcolor{red}{Exp}}\tau \sim \text{pr}\theta(G_{0,\tau}) = \sum_{\tau} \text{pr}\tau; \theta * G_{0,\tau}$$

The issue is that  $\text{pr}_{\tau; \theta}$  depends on

- the Environment's response at each step of  $\tau$ 
  - to the agent choosing actions  $\text{actseq}_{\text{tt}}$  in state  $\text{stateseq}_{\text{tt}}$  at step
- and the response is governed by probability  
$$\text{\transp}(\{ \text{stateseq}_{\text{tt+1}}, \text{rewseq}_{\text{tt+1}} | \text{state} = \text{stateseq}_{\text{tt}}, \text{act} = \text{actseq}_{\text{tt}} \})$$

BUT under the assumption of *Model-free* methods

- the Environment's Transition Probability is unknown

So,

- how can we compute the gradient of an expectation
- when don't know the probability distribution

# Policy Gradient Theorem

The *Policy Gradient Theorem* tells us how to compute the Gradient of the Expectation.

Most importantly

- the Environment's Transition probability *does not* appear
- so this results in an operational way to compute the Gradient needed for maximization of  $J(\theta)$

## Notes

- We simplify the presentation by assuming discount factor  $\gamma = 1$

We can write the Policy Gradient Theorem in two mathematically equivalent forms

### Episode Reward form

$$\nabla_{\theta} J(\theta) = \text{\textbackslash Exp} \tau \sim \pi_{\theta} \sum_{=0}^{|\tau|} \nabla_{\theta} \log \pi(\text{\textbackslash actseq}_{\tau,} | \text{\textbackslash stateseq}_{\tau,}) \text{\textbackslash rewseq}(\tau)$$

where

$$\text{\textbackslash rewseq}(\tau) = G_{0,\tau}$$

denotes the *episode reward*

This form is particularly useful

- when there is a *single reward* received at the end of the trajectory

### Notes

- To clarify that states, actions, rewards, returns, etc. depend on the specific trajectory  $\tau$ 
  - we add an extra subscript when necessary for clarification  
 $\text{\textbackslash rewseq}_{\tau,}$

## Periodic Reward form

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(\text{\color{red}\texttt{actseq}}_{\tau,t} | \text{\color{red}\texttt{stateseq}}_{\tau,t}) G_{\tau,t} \right]$$

where

$$G_{\tau,t}$$

is the return-to-go of the trajectory  $\tau$  from step (state  $\text{\color{red}\texttt{stateseq}}_{\tau,t}$ ).

This form is particularly useful

- when there are rewards at intermediate steps of the episode

Whichever way we write it

- the Policy Gradient Theorem is the foundation for all policy-based methods
- it tells us how to change the parameters  $\theta$  of the Policy NN in a direction leading to optimality

We will study a few of these methods in a later section.

# Computing $\nabla_{\theta} J(\theta)$

Our first step will be to turn the expectation

$$\langle \text{Exp} \rangle \tau \sim \text{\textbackslash prc} \tau \theta \text{\textbackslash rewseq}(\tau)$$

into a sum

$$\sum_{\tau \sim \text{\textbackslash prc} \tau \theta} \text{\textbackslash prc} \tau \theta * \text{\textbackslash rewseq}(\tau)$$

where

$$\text{\textbackslash prc} \tau \theta$$

is the probability of trajectory  $\tau$

## With stochastic policy and Environment

- for trajectory  $\tau$

$$\tau = \backslash \text{stateseq}_{\tau,0}, \backslash \text{actseq}_{\tau,0}, \backslash \text{rewseq}_{\tau,1}, \backslash \text{stateseq}_{\tau,1} \dots \backslash \text{stateseq}_{\tau,},$$
$$~~~~~\backslash \text{actseq}_{\tau,}, \backslash \text{rewseq}_{\tau,+1}, \backslash \text{stateseq}_{\tau,+1}, \dots$$

we can compute  $\backslash \text{prc}\tau\theta$

The presence of Environment Transition probability

$\text{\textbackslash transp}(\text{\textbackslash stateseq}_{\tau,+1}, \text{\textbackslash rewseq}_{\tau,+1} | \text{\textbackslash stateseq}_\tau, \text{\textbackslash actseq}_\tau)$

- is problematic
- as it is generally unknown
  - we would like the Model-free assumption to hold

Turning the expectation into a sum and taking the gradient of the expected Episode Reward

$$\begin{aligned}\nabla_{\theta} \sum_{\tau \sim \text{\textcolor{red}{prc}}\tau\theta} \text{\textcolor{red}{prc}}\tau\theta * \text{\textcolor{red}{rewseq}}(\tau) &= \sum_{\tau \sim \text{\textcolor{red}{prc}}\tau\theta} \nabla_{\theta} \text{\textcolor{red}{prc}}\tau\theta * \text{\textcolor{red}{rewseq}}(\tau) \\ &= \sum_{\tau \sim \text{\textcolor{red}{prc}}\tau\theta} (\text{\textcolor{red}{prc}}\tau\theta \nabla_{\theta} \log \text{\textcolor{red}{prc}}\tau\theta) \text{\textcolor{red}{rews}}\end{aligned}$$

---

We now substitute the chained probability previously derived for

$$\backslash \text{prc} \tau \theta$$

into the

$$\log \backslash \text{prc} \tau \theta$$

term above.

To summarize the proof:

$$\begin{aligned}\nabla_{\theta} \text{\textbackslash Exp} \tau \sim \text{\textbackslash pr} \theta \text{\textbackslash rewseq}(\tau) &= \nabla_{\theta} \sum_{\tau \sim \text{\textbackslash pr} \theta} \text{\textbackslash prc} \tau \theta * \text{\textbackslash rewseq}(\tau) \\ &= \sum_{\tau \sim \text{\textbackslash pr} \theta} (\text{\textbackslash prc} \tau \theta \nabla_{\theta} \log \text{\textbackslash prc} \tau \theta) * \text{\textbackslash rewseq}(\tau) \\ &= \sum_{\tau \sim \text{\textbackslash pr} \theta} \left( \text{\textbackslash prc} \tau \theta * \sum_{=0}^{|\tau|} \nabla_{\theta} \log \pi(\text{\textbackslash actseq}_{\tau, i} | \text{\textbackslash st}) \right) \\ &= \text{\textbackslash Exp} \tau \sim \text{\textbackslash prc} \tau \theta \sum_{=0}^{|\tau|} \nabla_{\theta} \log \pi(\text{\textbackslash actseq}_{\tau, i} | \text{\textbackslash st})\end{aligned}$$

---

## Key result

- We can evaluate the Gradient *without knowing* the model
  - i.e., Environment Transition Probability terms  $\backslash \text{transp}(\dots)$
- The expectation can be approximated by *sampling* trajectories
  - observe the *effect* of the Environment's distribution
  - *without knowing* it

**The convention is to write the expectation**

$\backslash \text{Exp}^\tau \sim \pi_\theta$  rather than  $\backslash \text{Exp}^\tau \sim \backslash \text{prc}^{\tau\theta}$

This is merely convention: they refer to the same distribution of episodes.

# Notes on Proof of Policy Gradient theorem

## Likelihood ratio trick

The likelihood ratio trick states that

- for a parameterized probability distribution  $p_\theta(x)$
- and a function  $f(x)$ :

the gradient of an expectation can be converted into an expectation over the gradient.

It is a simple consequence of the Derivative of a Log rule of calculus.

$$\begin{aligned}
\nabla_{\theta} \mathbb{E}_{x \sim p_{\theta}}[f(x)] &= \nabla_{\theta} \int p_{\theta}(x) f(x) dx \\
&= \int \nabla_{\theta} p_{\theta}(x) f(x) dx \\
&= \int f(x) \nabla_{\theta} p_{\theta}(x) dx \\
&= \int f(x) (p_{\theta}(x) \nabla_{\theta} \log p_{\theta}(x)) dx \\
&= \mathbb{E}_{x \sim p_{\theta}} [f(x) \nabla_{\theta} \log p_{\theta}(x)]
\end{aligned}$$

convert expectation to integral  
 move grad inside the integral  
 rearrange term  
 log trick:  
 $\nabla_{\theta} p_{\theta}(x) = p_{\theta}(x) \nabla_{\theta} \log p_{\theta}(x)$   
 convert integral back to expectation

The "log trick" follows from the rules of calculus

$$\nabla_{\theta} \log p_{\theta}(x) = \frac{1}{p_{\theta}(x)} * \nabla_{\theta} p_{\theta}(x) \quad \text{Calculus: grad of log, chain rule}$$

$$\nabla_{\theta} p_{\theta}(x) = p_{\theta}(x) * \nabla_{\theta} \log p_{\theta}(x) \quad \text{re-arranging terms}$$

**Why substituting  $\text{\texttt{rewseq}}_\tau$ , for  $G_{\tau,0}$  is algebraically the same**

Consider the two equivalent forms for expressing the Policy Gradient Theorem

$$\sum_{=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) \cdot R(\tau)$$

and

$$\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) \cdot G_t$$

where:

$$G = \sum_{k=0}^{T-1} \gamma^k r_k$$

and

$$R(\tau) = G_0$$

How can these two forms be mathematically equivalent ?

- since the first form involves  $G_0$ 
  - the rewards over all steps of the episode
- and the second form involves  $G_{\text{tt}}$ 
  - the future rewards from step onward
  - and  $G_{\text{tt}}$  and  $G'$  for ' $>$ ' include the same rewards

Algebraically they appear different.

The answer is that

- the two forms appear *within an expectation*
- which is evaluated over *future* time steps
- so the part of  $G_{\text{tt}}$  that reference *past rewards* is equal to 0 under the expectation

Here are the details:

- Step 1: Start with the total return  $R(\tau) = G_0$

Define  $L$  to be the expression for the first form of the Theorem:

$$L = \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot G_0 = G_0 \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

---

- Step 2: Expand  $G_0$  as the sum over rewards

$$G_0 = \sum_{k=0}^{T-1} \gamma^k r_k$$

Substitute into  $L$ :

$$L = \left( \sum_{k=0}^{T-1} \gamma^k r_k \right) \left( \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \right)$$

---

- Step 3: Express as a double sum

$$L = \sum_{t=0}^{T-1} \sum_{k=0}^{T-1} \gamma^k r_k \nabla_\theta \log \pi_\theta(a_t | s_t)$$


---

- Step 4: Separate sums over past and future rewards relative to  $t$

$$L = \sum_{t=0}^{T-1} \left( \sum_{k=0}^{t-1} \gamma^k r_k + \sum_{k=t}^{T-1} \gamma^k r_k \right) \nabla_\theta \log \pi_\theta(a_t | s_t)$$


---

- Step 5: Rewrite future rewards shifted by  $t$

Define  $j = k - t$ :

$$\sum_{k=t}^{T-1} \gamma^k r_k = \gamma^t \sum_{j=0}^{T-1-t} \gamma^j r_{t+j} = \gamma^t G_t$$


---

- Step 6: Substitute back into  $L$

$$L = \sum_{t=0}^{T-1} \left( \sum_{k=0}^{t-1} \gamma^k r_k + \gamma^t G_t \right) \nabla_\theta \log \pi_\theta(a_t | s_t)$$


---

- Step 7: Expectation zeroes out past rewards term

Because rewards before time  $t$  do not depend on action  $a_t$ , their expected contribution is zero:

$$\mathbb{E} \left[ \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot \sum_{k=0}^{t-1} \gamma^k r_k \right] = 0$$


---

- Step 8: Final form of the policy gradient

Thus,

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\pi} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot \gamma^t G_t \right]$$

which is the second form of expressing the Policy Gradient Theorem.

## Alternate Proof of the Policy Gradient Theorem (Per-Step Reward Perspective)

Let  $J(\theta)$  be the expected discounted sum of rewards under policy  $\pi_\theta$ :

$$J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^{T-1} \gamma^t r_t \right]$$

---

- Step 1: Expand the Expectation

Rewrite the expectation explicitly:

$$J(\theta) = \sum_{\tau} P_\theta(\tau) \left( \sum_{t=0}^{T-1} \gamma^t r_t \right)$$

---

- Step 2: Differentiation w.r.t.  $\theta$

$$\nabla_\theta J(\theta) = \sum_{\tau} \nabla_\theta P_\theta(\tau) \left( \sum_{t=0}^{T-1} \gamma^t r_t \right)$$

---

Apply the **likelihood ratio trick**:  $\nabla_\theta P_\theta(\tau) = P_\theta(\tau) \nabla_\theta \log P_\theta(\tau)$  So,

$$\nabla_\theta J(\theta) = \sum_{\tau} P_\theta(\tau) \nabla_\theta \log P_\theta(\tau) \left( \sum_{t=0}^{T-1} \gamma^t r_t \right)$$

Or equivalently,

$$\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[ \nabla_\theta \log P_\theta(\tau) \left( \sum_{t=0}^{T-1} \gamma^t r_t \right) \right]$$

- Step 3: Break Down  $\log P_\theta(\tau)$

Recall,  $\log P_\theta(\tau) = \sum_{t=0}^{T-1} \log \pi_\theta(a_t | s_t) + \text{terms independent of } \theta$  So,  
 $\nabla_\theta \log P_\theta(\tau) = \sum_{t'=0}^{T-1} \nabla_\theta \log \pi_\theta(a_{t'} | s_{t'})$  Substitute:

$$\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[ \sum_{t'=0}^{T-1} \nabla_\theta \log \pi_\theta(a_{t'} | s_{t'}) \cdot \left( \sum_{t=0}^{T-1} \gamma^t r_t \right) \right]$$


---

- Step 4: Swap Order of Summation

$$\nabla_\theta J(\theta) = \mathbb{E}_{\tau \sim \pi_\theta} \left[ \sum_{t'=0}^{T-1} \sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_{t'} | s_{t'}) \gamma^t r_t \right]$$

Switch the order:

$$= \mathbb{E}_{\tau \sim \pi_\theta} \left[ \sum_{t=0}^{T-1} \sum_{t'=0}^{T-1} \nabla_\theta \log \pi_\theta(a_{t'} | s_{t'}) \gamma^t r_t \right]$$


---

- Step 5: Analyze the causal relationship

The gradient w.r.t.  $a_{t'}$  can only affect rewards from  $t'$  onward (not earlier rewards due to the Markov property), so for  $t < t'$  the expectation is zero.

Thus, the only contributing terms are where  $t' \leq t$ :

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t'=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_{t'} | s_{t'}) \left( \sum_{t=t'}^{T-1} \gamma^t r_t \right) \right]$$

---

- Step 6: Recognize the return-to-go term

$$\sum_{t=t'}^{T-1} \gamma^t r_t = \gamma^{t'} \sum_{j=0}^{T-1-t'} \gamma^j r_{t'+j} = \gamma^{t'} G_{t'}$$

So, we may write:

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \cdot \gamma^t G_t \right]$$

Often,  $\gamma^t$  is absorbed into the definition of  $G_t$ .

---

- Step 7: Final form (policy gradient theorem)

$$\nabla_{\theta} J(\theta) = \mathbb{E}_{\tau \sim \pi_{\theta}} \left[ \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G_t \right]$$

---

### Conclusion:

By starting with the expectation of per-step rewards and applying the likelihood ratio trick, we arrive at the same policy gradient theorem:  
each action's gradient is weighted by the return-to-go from that time (not just the immediate reward).



# **Preview of Policy-Based Reinforcement Learning Methods**

We will subsequently present a number of Policy based methods.



# Actor-Critic

Value-based methods learn a function approximation of the *value* of a state or a state/action pair.

- policy is chosen based on the value of successor states

Simple Policy-based methods learn a parameterized policy function.

- using a NN to learn the policy
- using an objective function  $J(\theta)$  that depends on an approximation of either
  - the value  $\text{\statevalfun}(\text{\state})$  or  $G_{\text{tt}}$
  - or action/value function  $\text{\actvalfun}(\text{\state}, \text{\act})$

Actor-Critic-Policy-based methods used Neural Networks to learn

- *both* the value function and policy function approximations
- the agent is called the *Actor*
- the NN providing estimates of  $G_t$  or  $\text{\actvalfun}(\text{\state}, \text{\act})$  is called the *Critic*

- [RL tips and tricks \(\[https://stable-baselines3.readthedocs.io/en/master/guide/rl\\\_tips.html\]\(https://stable-baselines3.readthedocs.io/en/master/guide/rl\_tips.html\)\)](https://stable-baselines3.readthedocs.io/en/master/guide/rl_tips.html)
- [RL book contents \(<http://incompleteideas.net/book/RLbook2020.pdf#page=7>\)](http://incompleteideas.net/book/RLbook2020.pdf#page=7)
- [RL book notation \(<http://incompleteideas.net/book/RLbook2020.pdf#page=20>\)](http://incompleteideas.net/book/RLbook2020.pdf#page=20)

In [2]: `print("Done")`

Done

