

# Preferences vs Rewards

There are problems where providing *exact scalar rewards*

- is harder than *ranking* potential outputs.

For example

- I may prefer chocolate to vanilla
- but I can't quantify how much more

## Technically

- rewards form a total order
  - a reward has a magnitude
  - *all* rewards can be compared and ordered
- preferences form a partial order
  - we can order *some* pairs of outputs
  - without providing a magnitude

Good > Bad

Big > Small

Partial order:

Good > Small ?    undefined

Problems related to aligning the *style* of an LLM's output is a case of preferences.

- multiple answers may be "correct"
- but one answer may be "preferred"

For example

**Prompt:** "How do I change a tire?"

- **Reply A:** An accurate step-by-step answer.
- **Reply B:** A brief, incomplete answer.

Both replies are "correct" but the first is subjectively better.

An example of *Preference Data* is a triple

$$(x, y^+, y^-)$$

- input  $x$
- the preferred output  $y^+$
- the non-preferred output  $y^-$

## The case for preferences

Scenario	Why Preference Data?	Typical Example
RLHF & LLM alignment	Human feedback easier as comparisons	Choosing better LLM output
Hard-to-define or subjective “success”	Preference judgments more reliable	Dialogue, safety, style
Biased or noisy scalar rewards	Preferences less affected by outliers	Creative tasks, open-ended
Interpretability needs	Preferences can include rationales	Transparent value alignment
DPO-style methods	Direct optimization over preferences	Pairwise/choice based loss

In this module, we explore Reinforcement Learning for Preference Data.

## Example of a Preference Dataset

Here is a [link \(\[https://huggingface.co/datasets/HuggingFaceH4/ultrafeedback\\\_binarized\]\(https://huggingface.co/datasets/HuggingFaceH4/ultrafeedback\_binarized\)\)](https://huggingface.co/datasets/HuggingFaceH4/ultrafeedback_binarized) to the UltraFeedback dataset.

It is used to train an Assistant to be

- Helpful
  - answers the user's prompt; doesn't evade or decline
- Honest
  - gives a truthful answer

The methodology for constructing it is given [here](https://huggingface.co/datasets/HuggingFaceH4/ultrafeedback_binarized#dataset-card-for-ultrafeedback-binarized) ([https://huggingface.co/datasets/HuggingFaceH4/ultrafeedback\\_binarized#dataset-card-for-ultrafeedback-binarized](https://huggingface.co/datasets/HuggingFaceH4/ultrafeedback_binarized#dataset-card-for-ultrafeedback-binarized)).

- The authors gather a number of prompts across multiple domains.
- An AI assistant is asked to provide multiple responses to a prompt
- A second AI assistant
  - critiques the response based on, e.g., the Helpfulness criteria
  - gives a numerical evaluation
- Which creates a ranking of the responses to a prompt

## Synthetic Data in practice

The "binarized" dataset that we viewed

- selects the highest ranked answer as "Chosen"
- randomly selects the other responses as "Rejected"

Note the use of AI feedback rather than Human Feedback.

# **LLM Next Token prediction task and Reinforcement Learning**

One motivation for Preference Methods

- is to post-train an LLM
- to exhibit desirable characteristics
- which is often expressed with Preference Data
  - Preferred output vs. Non-Preferred output

We translate the LLM "Next Token Prediction" task to an instance of Reinforcement Learning

- in order to be able to use Reinforcement Learning for post-training an LLM

The Language Modeling task, formally, is

- Predict the Next Token conditional on the previously generated tokens of the response
- producing output  $y$  given input  $x$

For each output sequence  $\mathbf{y}$  of tokens

- We can associate a *state* corresponding to *each prefix* of  $y$ 
  - e.g,  $\text{stateseq} = \mathbf{y}_{[0:-1]}$
- An action is
  - choosing a token from the Vocabulary as the next output
- The policy  $\pi_\theta(\text{act} | \text{state})$  is thus equivalent to
$$\pi_\theta(\mathbf{y}_t | \mathbf{y}_{[0:-1]})$$
- the probability distribution of the next token, conditional on the prefix

Each sequence  $\mathbf{y}$  becomes an episode/trajectory.

We can write the probability  $\pi_\theta(\text{\textbackslash} \mathbf{y} | x)$

- of the *entire sequence*
- as the chained probabilities of each action given a state

This will be convenient

- in that we can compare the probabilities of
- a Preferred response  $y^+$  to a Non-Preferred response  $y^=$

rather than having to write the chained probability of each token in the sequences.

In [2]: `print("Done")`

Done

