

Value-based methods

The simplest model-based approach are *Value-based* methods.

They revolve around the idea of

- assigning a state value function $\text{\statevalfun}_\pi(\text{\state})$ to each state
 $\text{\state} \in \text{\States}$
$$\text{\statevalfun}_\pi : \text{\States} \rightarrow \text{\Reals}$$
- $\text{\statevalfun}_\pi(\text{\state})$ is an approximation of
$$\text{\E}_\pi(G \mid \text{\stateseq} = \text{\state})$$

the expected return achievable from state \state

Given statevalfun_π , the optimal deterministic policy is

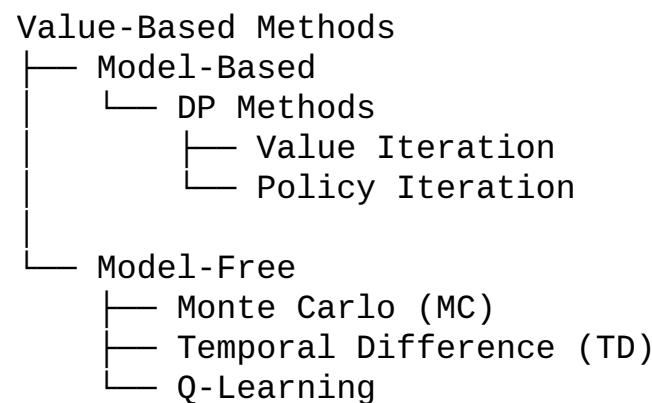
$$\begin{aligned} & \pi^*(\text{state}) \\ = & \underset{\text{act}}{\text{argmax}} \quad \text{act} \quad \text{statevalfun}_\pi(\text{state}') \\ & \text{transp}(\text{state}', \text{rew} | \text{state}, \text{act}) \neq 0 \end{aligned}$$

- From state state
- Choose the action act
- that results in next state state'
- with maximal $\text{statevalfun}_\pi(\text{state}')$

Note: the argmax results in a *deterministic* policy.

We will show two broad classes of Value-based methods

- Model-based
- Model-free



Model-based Value methods

Using a value-based method is practical only if we can discover the state value function $\text{\textbackslash statevalfun}$, which is initially unknown.

Most methods are iterative in nature and create a sequence of increasingly accurate approximations of $\text{\textbackslash statevalfun}$

$$\text{\textbackslash statevalfun}_{\pi,0} \dots \text{\textbackslash statevalfun}_{\pi,k} \dots$$

In the limit

$$\lim_{k \rightarrow \infty} \text{\textbackslash statevalfun}_{\pi,k} = \text{\textbackslash statevalfun}_{\pi}$$

As we obtain an improved approximation $\text{\textbackslash statevalfun}_{\pi,k+1}$

- we may reflect this improved knowledge by updating the policy

Thus, we periodically improve the policy based on improved approximations of
 $\text{\textbackslash statevalfun}_\pi$

This results in a sequence of increasingly accurate approximations of the policy π

$$\pi_0, \dots, \pi_p, \dots$$

which hopefully converges to π^* .

Dynamic Programming

Given a model that describes the behavior of the Environment

- we can determine the value state function via Dynamic Programming (DP) based techniques.

as follows.

The expected returns from state \state

$$\mathbb{E}_\pi(G_1 | \text{stateseq} = \text{state})$$

are computed via the *Bellman Equation*

$$\begin{aligned}\text{\statevalfun}_\pi(\text{stateseq}) &= \mathbb{E}_\pi(G_1 | \text{stateseq} = \text{state}) \\ &= \mathbb{E}_\pi(\text{rewseq}_{+1} + \text{disc}G_{+1} | \text{stateseq} = \text{state}) \\ &= \mathbb{E}_\pi(\text{rewseq}_{+1} + \text{disc}\text{\statevalfun}_\pi(\text{stateseq})\end{aligned}$$

The Bellman Equation asserts that

$$\text{statevalfun}_\pi(\text{stateseq})$$

can be derived from

- the immediate reward rewseq_{+1}
- and the discounted (by γ) value of the successor state stateseq_{+1}
 $\text{statevalfun}_\pi(\text{stateseq}_{+1})$

This recursive equation terminates

- since the successor state stateseq_{+1} on the RHS
- is one transition closer to the end of episode π
- than the LHS state stateseq

Interpreting the Expectation

Expanding the expectation into a sum

$$\text{\statevalfun}_\pi(\text{\stateseq}) = \sum_{\text{\act}} \pi(\text{\act}, \text{\stateseq}) \sum_{\text{\state}', \text{\rew}} \text{\transp}(\text{\state}', \text{\rew} | \text{\stateseq}, \text{\act}) (\text{\rew} + \text{\disc} \text{\statevalf}$$



This equation guides our approximation of $\text{statevalfun}_\pi(\text{\stateseq})$, given current policy π

- for each action \act that can be chosen by the Agent in state \stateseq_tt
 - with probability $\pi(\text{\act}, \text{\stateseq})$ (stochastic policy)
- use the return received from taking the action
 - determined by the Environment, which chooses
 - immediate reward \rew
 - successor state $\text{\state}'$
 - with probability $\text{transp}(\text{\state}', \text{\rew} | \text{\stateseq}, \text{\act})$ (stochastic environment)
- use the current approximation v of the value of the successor state $\text{\state}'$

This equation can be computed only if we know the behavior of the Environment

- the reward and successor state chosen by the Environment
- given the Agent choosing action act in state \stateseq_tt
- i.e.

$$\text{\transp}(\text{\state}', \text{\rew} | \text{\stateseq}, \text{\act})$$

Hence: this is *model-based* and not model-free.

Model-based methods have access to the Environment's behavior.

This may come about because

- the model is given to us
- we incorporate methods to *learn* a model simultaneously with learning the Value function

For now: we assume the model is given to us.

Simplification for deterministic policy or environment

For deterministic policy π

- as given by

$$\pi(\text{\state})$$

$$= \operatorname{argmax}_{\text{\act}} \quad \text{\act} \\ \text{\transp}(\text{\state}', \text{\rew} | \text{\state}, \text{\act}) \neq 0$$

- chose the action resulting in successor state with maximum state value

all the probability is concentrated at single action $\boxed{\text{\actseq_tt}}$

$$\pi(\text{\stateseq}, \text{\act}^*) = 1 \text{ for } \text{\act}^*$$

$$= \operatorname{argmax}_{\text{\act}} \quad \text{\act} \\ \text{\transp}(\text{\state}', \text{\rew} | \text{\state}, \text{\act}) \neq 0$$

so you can drop the $\sum_{\text{\act}} \pi(\text{\act}, \text{\stateseq})$ from the equation

- and equate $\boxed{\text{\act} = \text{\actseq_tt}}$

When the Environment is deterministic all the probability is concentrated at a single response:

$$\text{transp}(\{ \text{stateseq}_{\{\text{tt}+1\}}, \text{rewseq}_{\{\text{tt}+1\}} | \text{stateseq}_{\text{tt}}, \text{actseq}_{\text{tt}} \}) = 1$$

and the Agent's policy chooses the $\text{actseq}_{\text{tt}}$ that results in a successor state

$\text{state}' = \text{stateseq}_{+1}$ that maximizes $\text{statevalfun}(\text{state}')$

$$\text{statevalfun}_\pi(\text{stateseq}_{+1}) = \max \text{state}' \text{statevalfun}_\pi(\text{state}')$$

Thus, in the case of deterministic Policy and Environment, the RHS becomes

$$\text{disc} (\text{rewseq}_{+1} + \max \text{state}' \text{statevalfun}_\pi(\text{state}'))$$

which is how it typically appears.

Value iteration method

The simplest method is to

- iteratively update the Value function
 - until convergence
- derive the *final* Policy from the Value function
 - chose the action leading to highest return
 - based on the Value function

Pseudo code for Value Iteration

Here is some pseudo-code

Value iteration

Value Iteration, for estimating $\pi \approx \pi_*$

Algorithm parameter: a small threshold $\theta > 0$ determining accuracy of estimation
Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop:

```
| Δ ← 0
| Loop for each  $s \in \mathcal{S}$ :
|    $v \leftarrow V(s)$ 
|    $V(s) \leftarrow \max_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$ 
|   Δ ← max(Δ, |v - V(s)|)
until Δ < θ
```

Output a deterministic policy, $\pi \approx \pi_*$, such that
 $\pi(s) = \arg \max_a \sum_{s',r} p(s',r|s,a) [r + \gamma V(s')]$

Attribution: <http://incompleteideas.net/book/RLbook2020.pdf#page=105>
[\(http://incompleteideas.net/book/RLbook2020.pdf#page=105\)](http://incompleteideas.net/book/RLbook2020.pdf#page=105)

Value iteration

Initialize value function $V(s)$ arbitrarily (e.g., zero for all states)
Repeat:
delta = 0
For each state s :
old_value = $V(s)$
 $V(s) = \max_{a} [R(s, a) + \gamma * \sum_{s'} [P(s' | s, a) * V(s')]]$
delta = max(delta, |old_value - V(s)|)
Until delta < threshold
Derive policy after value function converges
For each state s : $\pi(s) = \arg\max_a [R(s, a) + \gamma * \sum_{s'} [P(s' | s, a) * V(s')]]$
Return π, V

Subtlety

Notice that the Bellman equation in the code is modified

- to reflect the deterministic, optimal choice of action

max over a { ... }

rather than an expectation over all possible actions

$$\sum_{\backslash \text{act}} \pi(\backslash \text{act}, \backslash \text{stateseq}, \{ \dots \})$$

Policy iteration method

Rather than updating the Policy once (after Value function convergence)

- we introduce a method that periodically updates the Policy.

Policy iteration is an algorithm that improves π_p to π_{p+1} by alternating two steps during round p

The algorithm alternates between

- Policy evaluation
 - update the estimate of $\text{\textbackslash statevalfun}_{\pi_p}$ to $\text{\textbackslash statevalfun}_{\pi_{p+1}}$
- Policy improvement
 - update π_p to π_{p+1} using the newly updated $\text{\textbackslash statevalfun}_{\pi_{p+1}}$

Here is some pseudo-code

Policy iteration

Policy Iteration (using iterative policy evaluation) for estimating $\pi \approx \pi_*$

1. Initialization

$V(s) \in \mathbb{R}$ and $\pi(s) \in \mathcal{A}(s)$ arbitrarily for all $s \in \mathcal{S}$; $V(\text{terminal}) \doteq 0$

2. Policy Evaluation

Loop:

$$\Delta \leftarrow 0$$

Loop for each $s \in \mathcal{S}$:

$$v \leftarrow V(s)$$

$$V(s) \leftarrow \sum_{s',r} p(s', r | s, \pi(s)) [r + \gamma V(s')]$$

$$\Delta \leftarrow \max(\Delta, |v - V(s)|)$$

until $\Delta < \theta$ (a small positive number determining the accuracy of estimation)

3. Policy Improvement

policy-stable \leftarrow true

For each $s \in \mathcal{S}$:

$$\textit{old-action} \leftarrow \pi(s)$$

$$\pi(s) \leftarrow \operatorname{argmax}_a \sum_{s',r} p(s', r | s, a) [r + \gamma V(s')]$$

If $\textit{old-action} \neq \pi(s)$, then *policy-stable* \leftarrow false

If *policy-stable*, then stop and return $V \approx v_*$ and $\pi \approx \pi_*$; else go to 2

Attribution: <http://incompleteideas.net/book/RLbook2020.pdf#page=102>
[\(http://incompleteideas.net/book/RLbook2020.pdf#page=102\)](http://incompleteideas.net/book/RLbook2020.pdf#page=102).

Pseudo code for Policy Iteration

Policy iteration

Initialize policy π arbitrarily (e.g., random policy) Initialize value function $V(s)$ arbitrarily (e.g., zero for all states)
Repeat: # Policy Evaluation (compute V for current policy π)
Repeat: For each state s : $V(s) = R(s, \pi(s)) + \gamma * \text{sum over } s' [P(s' | s, \pi(s)) * V(s')]$ Until $V(s)$ converges (changes smaller than threshold)
Policy Improvement (update policy based on current V)
 $\text{policy_stable} = \text{True}$
For each state s : $\text{old_action} = \pi(s)$ $\pi(s) = \text{argmax over } a [R(s, a) + \gamma * \text{sum over } s' [P(s' | s, a) * V(s')]]$
if $\text{old_action} \neq \pi(s)$: $\text{policy_stable} = \text{False}$
Until policy_stable is True
Return π, V

Iterative improvement of $\text{\textbackslash statevalfun}_{\pi,k}$ to $\text{\textbackslash statevalfun}_{\pi,k+1}$ is via the equation

$$\text{\textbackslash statevalfun}_{\pi,k+1}(\text{\textbackslash state}) = \sum_{\text{\textbackslash state}', \text{\textbackslash rew}} \text{\textbackslash transp}(\text{\textbackslash state}', r | \text{\textbackslash state}, \pi(\text{\textbackslash state}))$$

We continue iterating until, for all states $\text{\textbackslash state}$,

- the difference between $\text{\textbackslash statevalfun}_{\pi,k+1}(\text{\textbackslash state})$ and $\text{\textbackslash statevalfun}_{\pi,k}(\text{\textbackslash state})$ is smaller than a threshold value θ .
-

Subtleties

The Policy Evaluation equation

$$\text{\textbackslash statevalfun}_{\pi,k+1}(\text{\textbackslash state}) = \sum_{\text{\textbackslash state}', \text{\textbackslash rew}} \text{\textbackslash transp}(\text{\textbackslash state}', r | \text{\textbackslash state}, \pi(\text{\textbackslash state}))$$

is *similar* (but not identical) to the Bellman equation.

- we assume a *deterministic* Policy
 - $\pi(\text{\textbackslash state})$ in the pseudo-code is a single choice
 - i.e., the *optimal* one
 - choosing the action that leads to the successor state $\text{\textbackslash state}'$ with maximum value
 - as determined by the current Value function
-

The deterministic policy π is improved from π_p to π_{p+1} with the Policy improvement step

$$\pi'_{p+1}(\text{\textbackslash state}) = \underset{\sum_{\text{\textbackslash state}', \text{\textbackslash rew}} \text{\textbackslash transp}(\text{\textbackslash state}', \text{\textbackslash rew} | \text{\textbackslash state}, \text{\textbackslash act})(\text{\textbackslash rew} + \text{\textbackslash disc})}{\text{\textbackslash argmax} \text{\textbackslash act}}$$

That is: the agent in state $\text{\textbackslash state}$ chooses the action with maximal return.

The alternation between Policy Evaluation and Policy Improvement

- continues until the updated and previous policy are identical
-

Convergence to the correct value function is guaranteed

- Each update transfers information to a state from all successor states
- This ensures that information about all states eventually reaches each affected state

The updated policy is no worse than the previous one.

- Reference: *Policy Improvement Theorem*

There are a finite number of *deterministic* policies

$$|\setminus \text{Actions}|^{\setminus |\text{States}|}$$

So Policy Iteration eventually arrives at the optimal policy.

The advantage of alternating between Policy Evaluation and Policy Improvement

- faster convergence
 - the Value function under the current policy is fully evaluated
 - before the Policy is updated
- more stable convergence
 - Policy doesn't change until Value function (under current policy) is fully known

Key Differences Between Value Iteration and Policy Iteration

Feature	Value Iteration	Policy Iteration
Approach	Updates value function iteratively using the Bellman Optimality Equation in one step	Alternates between full policy evaluation and policy improvement steps
Policy Handling	Policy is implicitly updated after value convergence	Explicitly maintains and updates policy each iteration
Initialization	Starts with an initial value function	Starts with an initial policy
Iteration Steps	Single step combining evaluation and improvement	Two-step process: separate evaluation and improvement
Convergence Criterion	Value function converges (changes below a threshold)	Policy stabilizes (no change between iterations)
Computation per Iteration	Potentially more expensive per iteration (max over all actions for each state)	More computationally intensive due to full policy evaluation, but fewer total iterations needed
Number of Iterations	Typically more iterations	Usually fewer iterations
Complexity	Simpler to implement	More complex implementation
Suitability	Suitable for smaller state spaces or when full policy evaluation is expensive	Can handle larger state spaces more efficiently when full evaluation is feasible
Policy Updates	Policy derived after convergence of value function	Policy updated after each evaluation phase

Here is a comparison of the Value and Policy Iteration methods.

Feature	Value Iteration	Policy Iteration
Approach	Updates value function until convergence	Alternates between value evaluation and improvement
Convergence	When value function $V(s)$ stabilizes	When policy $\pi(s)$ stops improving
Computational Cost	Higher per iteration, simpler logic	Lower per iteration, more complex
Speed	Requires more iterations	Fewer iterations; often faster
Policy Output	Extracted after value convergence	Updated during each iteration

Finding the best action in a Value-based method

The Value-based methods don't directly give you a policy

- the Value function gives you the best successor state $\text{\textbackslash state}'$ from current state $\text{\textbackslash state}$
- but it doesn't directly tell you the action $\text{\textbackslash act}$ that leads to $\text{\textbackslash state}'$

In order to find $\text{\textbackslash act}$ you either

- need a model
 - search over all possible actions, using the model to determine the value of action's successor state
- use actual experience to estimate the effect of performing action $\text{\textbackslash act}$ in state $\text{\textbackslash state}$

Here is some pseudo-code that uses a model to determine the best action:

```
# For each possible action a in current state s:  
for each action a in actions:  
    # Take action a from state s in the environment  
    observe next state s', reward r  
    # Estimate value for taking action a from s  
    A[a] = r + gamma * V[s']  
# Find the maximum estimated action value  
A_max = max over a of A[a]  
  
# Update the value function for state s  
V[s] = V[s] + alpha * (A_max - V[s])
```

Value-based methods: Advantages/Disadvantages

Credit assignment (implied intermediate rewards)

In many episodes,

$$\dots \backslash \text{stateseq}, \backslash \text{actseq}, \backslash \text{rewseq}_{+1}, \dots$$

the only reward comes from entering the terminal state, thus

$$\backslash \text{rewseq}_{+1} = 0$$

for many time steps .

With $\backslash \text{statevalfun}_\pi$ in-hand

- we can interpret the *increment* in value
 $\backslash \text{statevalfun}_\pi(\backslash \text{stateseq}_{+1}) - \backslash \text{statevalfun}_\pi(\backslash \text{stateseq})$
of the action that takes us from $\boxed{\backslash \text{stateseq_tt}}$ to $\backslash \text{stateseq}_{+1}$
- as an implicit reward that provides immediate feedback

Limitations: deterministic policy; discrete actions

Policy is deterministic; can't have stochastic policy

$$\pi^*(\text{\textbackslash state}) = \text{\textbackslash argmax}\text{\textbackslash act}\text{\textbackslash actvalfun}(\text{\textbackslash state}, \text{\textbackslash act})$$

- actions are discrete, not continuous
 - the magnitude of angles (when turning) or velocity (when moving) are not continuous
 - a consequence of the max $\text{\textbackslash act}'$

Model-free methods

In the absence of a model

- We will learn the dynamics of the Environment
- Maintaining *estimates* of the optimal policy
- Updating the estimates through the feedback of reward/next state transition

Greed is not (always) good

Before presenting specific methods, we discuss the concepts of exploitation and exploration.

- which will be used in subsequent methods

Given the current approximation of the action value function

$\text{\textbackslash actvalfun}_\pi(\text{\textbackslash state}, \text{\textbackslash act})$

- the obvious policy choice for action $\boxed{\text{\textbackslash actseq_}\text{\texttt{t}}}$ is the one with
 $\max a' \text{\textbackslash actvalfun}_\pi(\text{\textbackslash state}, a')$

The problem with this greedy choice of action is that, initially, our estimate of the true $\text{\textbackslash actvalfun}_\pi$ is inaccurate.

- by choosing the current estimate of "optimal" action
- we may fail to ever choose the true optimal
- and we will never learn the optimal action as a result

Choosing the current "best" action is called *exploitation*.

Sometimes *exploration* (making a seemingly sub-optimal choice) sacrifices short term gain for long term gain.

This is called the *exploration-exploitation* trade-off.

Note

The example above

- illustrates the concept using the Q-learning method (to be introduced subsequently)
- but would be similar for any other method that updates and estimate

This reflects the primary difference between Model-based and Model-free methods

- Model-based methods have knowledge of *all* possible transitions
- Model-free methods have knowledge only of the single transition reflected by the chosen action

The single experience sampled by Model-free methods is *noisy*

- hence, we moderate changes to estimates

Update Style	Method	Learning Rate?	Basis of Update
Replacement	Model-based (DP)	No	All possible transitions
Incremental	Model-free (TD, Q, MC)	Yes (α)	One sampled transition

Temporal Difference (TD) methods

There is a family of methods (TD) in which

- there is some estimate associated with a state \stateseq_tt
 - $\text{\statevalfun}_{\pi,k}(\text{\stateseq})$ for V-learning
 - $\text{\actvalfun}_{\pi,k}(\text{\stateseq}, \text{\actseq})$ for Q-learning
- the estimate evolves sequentially
- by adding an *increment* to the current estimate
 - e.g.,
$$\text{\statevalfun}_{\pi,k+1}(\text{\stateseq}) = \text{\statevalfun}_{\pi,k}(\text{\stateseq}) + \delta$$

δ is called the *Temporal Difference Error*

The increment is often moderated by a *learning rate* α

$$\text{\statevalfun}_{\pi,k+1}(\text{\stateseq}) = \text{\statevalfun}_{\pi,k}(\text{\stateseq}) + \alpha * \delta$$

How do we obtain the increment δ ?

- via the Bellman-like equation that defines the estimate

$$\begin{aligned}\text{\color{red} \statevalfun}_{\pi,k+1}(\text{\color{red} \stateseq}) &= \text{\color{red} \rewseq}_{+1} \\ &\quad + \\ &\quad \text{\color{red} \disc \statevalfun}_{\pi,k}(\text{\color{red} \stateseq}_{+1}) \\ \delta &= \text{\color{red} \statevalfun}_{\pi,k+1}(\text{\color{red} \stateseq}) - \text{\color{red} \statevalfun}_{\pi,k}(\text{\color{red} \stateseq})\end{aligned}$$

Backups

This equation utilizes a common technique referred to as *value backup*

- the value of successor states
 $\text{\statevalfun}_{\pi,k}(\text{\stateseq}_{+1})$
- are propagated "back" to the current state $\boxed{\text{\stateseq_}\text{\tt}}$

The term

$$\text{\rewseq}_{+1} + \text{\disc}\text{\statevalfun}_{\pi,k}(\text{\stateseq}_{+1})$$

is called the *target*.

The target is considered a "more informed" estimate of true

$$\text{\textcolor{red}{statevalfun}}_{\pi}(\text{\textcolor{red}{stateseq}})$$

than the current estimate $\text{\textcolor{red}{statevalfun}}_{\pi,k}(\text{\textcolor{red}{stateseq}})$ as it includes information about

- the reward $\text{\textcolor{red}{rewseq}}_{+1}$
 - which will only be received *after* we take the action $\boxed{\text{\textcolor{black}{actseq}} \text{\textcolor{black}{\tt}}}$
- and a successor state: $\text{\textcolor{red}{statevalfun}}_{\pi,k}(\text{\textcolor{red}{stateseq}}_{+1})$

We improve our estimate of $\text{\textcolor{red}{statevalfun}}_{\pi}(\text{\textcolor{red}{stateseq}})$ to

$$\text{\textcolor{red}{statevalfun}}_{\pi,k+1}(\text{\textcolor{red}{stateseq}})$$

- by moving in the direction of the target

Bootstrapping

Note that δ depends on an *estimated* value

- $\text{\statevalfun}_{\pi,k}(\text{\stateseq}_{+1})$

So our new estimate for $\text{\statevalfun}_{\pi,k+1}(\text{\stateseq})$ is based on estimate $\text{\statevalfun}_{\pi,k}(\text{\stateseq}_{+1})$.

When one estimate is based on another estimate

- this is called *bootstrapping*

Bootstrapping is the *defining characteristic* of the TD technique.

p-step ahead updates

The update above is based on the immediate (one step ahead) reward.

This method is called *1-step TD*

We can generalize this to *p-step TD* which uses

- rewards $\{\text{\textbackslash rewseq}_t$

| $<'\leq$

$+p\}$

- $\text{\textbackslash statevalfun}_{\pi,k}$

$(\text{\textbackslash stateseq}_{+p})$

$$\delta = \left[\sum_{i=1}^p \gamma^{i-1} R_{t+i} + \gamma^p \text{\textbackslash statevalfun}_{\pi,k}(\text{\textbackslash stateseq}_{t+p}) \right] - \text{\textbackslash statevalfun}_{\pi,k}(\text{\textbackslash stateseq})$$

Note that the value used for the end state $\text{\textbackslash stateseq}_{+p}$ is

- the pre-update value
 $\text{\textbackslash statevalfun}_{\pi,k}(\text{\textbackslash stateseq}_{+p})$
- rather than the post-update value
 $\text{\textbackslash statevalfun}_{\pi,k+1}(\text{\textbackslash stateseq}_{+p})$

That is

- all backups occur simultaneously
- not sequentially
 - so the update to $\text{\textbackslash statevalfun}_\pi(\text{\textbackslash stateseq}_{+1})$ does not influence the update to $\text{\textbackslash statevalfun}_\pi(\text{\textbackslash stateseq})$
 - even though we might happen to evaluate it first (bottom up)

To illustrate the backups:

State_0 --a--> State_1 --a--> ... --a--> State_T (terminal)

TD: Update $V(s)$ at every transition using $V(s_{\{t+1\}})$

Pseudo-code for p-step TD

We give some code below.

The apparent complexity of the code arises

- because p-step TD, for $p > 1$
- involves returns that won't be experienced
- until $(p - 1)$ steps after step

So the update for $\text{\textbackslash statevalfun}_{\pi,k}(\text{\textbackslash stateseq})$ is delayed $p - 1$ steps

- until we accumulate the remaining rewards
- at the beginning of update ($k + 1$) we already have the value of $\text{\textbackslash statevalfun}_{\pi}(\text{\textbackslash stateseq}_{+p})$ needed
 - i.e., $\text{\textbackslash statevalfun}_{\pi,k}(\text{\textbackslash stateseq}_{+p})$

```
from collections import defaultdict

gamma = 0.99           # Discount factor
alpha = 0.1            # Learning rate
k = 3                 # Number of steps (set k as needed)
V = defaultdict(float) # State-value function

def td_k_episode(env, policy, k):
    states = []
    rewards = []

    s = env.reset()
    states.append(s)
    done = False
    t = 0

    while not done:
        a = policy(s)
        s_next, r, done, info = env.step(a)
        rewards.append(r)
```

Q-learning: Action-value function

We present a TD method called Q-learning.

The Value-function method makes it a little awkward to extract the action leading to the successor state with highest value.

- The state-value function associates the return (discounted future rewards) with a state
- without explicit reference to the action that leads to this return
- to find the best action, we either
 - need a model; use it to measure the value of each action
 - use experience to simulate the effect of each action

A simple extension of the Value function into an *Action-Value function* simplifies the determination of the next action:

$$\text{\textcolor{red}{actvalfun}}_{\pi} : \text{\textcolor{red}{state}} \times \text{\textcolor{red}{act}} \rightarrow \text{\textcolor{red}{Reals}}$$

The Action-Value function maps a state and chosen action into the value of the successor state.

So the action act^* that leads to maximum $\text{statevalfun}_\pi(\text{state}')$ is easily obtained from the Action-Value function via

$$\pi^*(\text{state}) = \text{argmax}_{\text{act}} \text{actvalfun}(\text{state}, \text{act})$$

Note: the argmax results in a *deterministic* policy just as before

The Bellman equation for the Action-Value function is

$$\text{\textbackslash actvalfun}_\pi(\text{\textbackslash state}, \text{\textbackslash act}) = \text{\textbackslash E}_\pi (\text{\textbackslash rewseq}_{+1} + \text{\textbackslash disc max } a' \text{\textbackslash actvalfun}_\pi(\text{\textbackslash st}$$

Thus

$$\text{\textbackslash actvalfun}_\pi(\text{\textbackslash state}, \text{\textbackslash act}) = \text{\textbackslash E}_\pi(G | \text{\textbackslash stateseq_state}, \text{\textbackslash actseq_a})$$

Q-learning is a method of learning the Action-Value function.

It implements the actvalfun_π function (mapping state/action pairs to return) via a *tabular* lookup

- table is built dynamically through experience

It updates the estimated actvalfun_π similar to the method used in the Model-based approach.

- except that it can only update the action *taken* by the current policy
- the experience is gathered by taking the action

Pseudo code for Q-learning

```
Initialize Q(s, a) arbitrarily for all states s and actions a (often Q(s, a) = 0)

Set learning rate alpha, discount factor gamma, exploration rate epsilon

For episode = 1 to number_of_episodes:
    Initialize state s

    Repeat until s is terminal:
        With probability epsilon:
            Choose a random action a (exploration)
        Otherwise:
            Choose action a = argmax_a Q(s, a) (exploitation)

        Take action a, observe reward r and next state s'

        Update Q(s, a) using:
            Q(s, a) = Q(s, a) + alpha * [r + gamma * max_{a'} Q(s', a') - Q(s, a)]
```

In the above code

- we build a table Q for the action value function $\text{\textbackslash actvalfun}_\pi$
- You can see the ϵ -greedy strategy

With probability epsilon:

Choose a random action a (exploration)

Otherwise:

Choose action $a = \text{argmax}_a Q(s, a)$ (exploitation)

- Update assumes subsequent action choices are greedy

$\text{max}_{\{a'\}} Q(s', a')$

- Model-free: Notice that we don't make use of
 - $\text{\textbackslash transp}(\text{\textbackslash state}', \text{\textbackslash rew} | \text{\textbackslash state}, \text{\textbackslash act})$
 - or any reward other than the one received by taking the chosen action

An ϵ -greedy policy manages the exploration-exploitation trade-off

- by choosing the (current) best action with probability $(1 - \epsilon)$
- choosing a random action with probability ϵ

Q-learning will use such an ϵ -greedy policy .

Deep Q-learning (DQN)

The method for Q-learning presented involved creating a *table* implementing the mapping actvalfun_π .

This is only practical when the size of the table is small

- the number of states for many problems (e.g., games) is extremely large
- not practical

Deep Q-Learning

- treats actvalfun_π as a function
- which is approximated by a Neural Network (the *Deep Q-Network (DQN)*)

The function is *trained*

- to reproduce the value calculated in ordinary Q-learning
- via a Loss function (MSE)
 - that minimizes the difference between
 - the NN output
 - and the mathematical definition of the Q function.

That is:

- We train the NN function $Q_\theta(\text{\stateseq}, \text{\actseq})$ to approximate true value $\text{\actvalfun}_\pi(\text{\stateseq}, \text{\actseq})$
- Using an MSE per-example loss
$$(Q_\theta(\text{\stateseq}, \text{\actseq}) - \text{\actvalfun}_\pi(\text{\stateseq}, \text{\actseq}))^2$$

The loss is

- calculated on a mini-batch of steps
- gradient descent on the batch loss results in an update to the NN parameters Θ

So one fundamental addition to the ordinary Q-learning algorithm

- record state transitions in a Replay Memory

Mini-batches are sampled from the Replay Memory

The Replay Memory enables *Experience Replay*

- off-line training
 - can train on the *same state transition* without executing the episode again
 - efficient use of transitions (can reuse)
- batch creation
- having previous transitions in a batch
 - prevents forgetting that might occur by only seeing *new* transitions

Pseudo code for Deep Q-learning

```
Initialize replay memory D to capacity N
Initialize main Q-network with random weights  $\theta$ 
Initialize target Q-network with weights  $\theta^- = \theta$ 

Set exploration rate  $\varepsilon$ , discount factor  $\gamma$ , learning rate  $\alpha$ 

For episode = 1 to M:
    Initialize state s

    For each step in episode:
        With probability  $\varepsilon$ :
            Choose random action a (exploration)
        Else:
            Choose  $a = \text{argmax}_a Q(s, a; \theta)$  (exploitation)

        Take action a, observe reward r and next state  $s'$ 

        Store transition  $(s, a, r, s')$  in replay memory D

        Sample random mini-batch of transitions  $(s_j, a_j, r_j, s'_j)$  from
```

Key aspects of the above code:

- ϵ -greedy choice of action
- Replay memory D
 - stores experiences *using the policy that was in effect when the experience was created*
 - not necessarily the current policy (different values of parameter θ)
 - a batch of experiences is sampled from D to train the NN computing Q_pred

- Updates in a *batch*
 - in the "basic" value-action method
 - $\backslash \text{actvalfun}_\pi$ is updated for *each action* of the Agent
 - in Q-learning
 - Q (analogous to $\backslash \text{actvalfun}_\pi$) is updated with *multiple, randomly chosen prior* actions of the Agent
 - mini-batch: target computed for each example in the batch


```
For each sample in the batch:
      y_j = r_j + γ * max_{a'} Q(s'_j, a'; θ^-)    # Target Q-value
      # Predicted Q-value
      Q_pred = Q(s_j, a_j; θ)
```
 - per-example loss: compare target y_j with prediction Q_{pred}

- Loss function: MSE between Q_{pred} and calculated y_j

Compute loss = mean squared error between y_j and Q_{pred}

- Q_{pred} is the value predicted by the NN
 - for a *batch* of examples
- y_j is the "target value":
 - for the batch
 - the true value that the NN is trying to match
 - *defined* by the same calculation as regular Q-learning (Bellman equation)

$$r + \gamma \max_{a'} Q(s', a')$$
 but $Q(s', a')$ replaced by NN calculated $Q(s', a'; \theta^-)$
 - where θ^- are the NN's lagged weights

The advantage of updating on multiple actions rather than just the current one

- smoother updates
 - changes "averaged" over many actions, not just the current one
- avoids catastrophic forgetting
 - emphasizes retention of past knowledge, not just current action

This is all similar to the reason we use Mini-Batch Gradient Descent in Neural Networks.

Why the lagged weights θ^- in computing the target value via $Q(s', a'; \theta^-)$?

To introduce stability in training.

If we don't lag the weights: the targets computed for other mini-batches will be based on different weights

- so we have a moving target as well as a moving function

The lagged weights are periodically synchronized with the most recent weights.

Q-learning (Action-value function) vs V-learning (Value function)

The main difference between the Value function and the Action-Value function

- Value function: the value of a state is averaged over *all* actions
- Action-Value function
 - absent a model: estimate effect of a single action

The averaging in V-learning can make the estimates

- more accurate and less noisy
- particularly when the action chosen is via the "max"

(We will quantify this in the section on Bias and Variance)

Monte Carlo (MC) method (not a TD method)

In the *Monte Carlo* method

- the update δ to *each state*
- is based on the return to go from the state
 - return accumulated *to the end of the episode*
- rather than the return of just the next p experiences

Loosely, it is TD(∞)

Note

Technically: MC is *not* a TD method

- as it does not rely on bootstrapping
 - not based on other estimated values

To highlight the difference between TD and MC:

- consider an episode from initial $\text{\textbackslash stateseq}_1$ to $\text{\textbackslash stateseq}_2$ to terminal state $\text{\textbackslash stateseq}_3$
- earning a reward of +1 on the transition from $\text{\textbackslash stateseq}_1$ and $\text{\textbackslash stateseq}_2$
- with initial state values

$$\text{\textbackslash statevalfun}_{\pi}(\text{\textbackslash stateseq}_{1,0}) = \text{\textbackslash statevalfun}_{\pi}(\text{\textbackslash stateseq}_{2,0}) = 0$$

After episode 1

- using TD

$$\text{\textbackslash statevalfun}_{\pi}(\text{\textbackslash stateseq}_{1,1}) = \text{\textbackslash statevalfun}_{\pi}(\text{\textbackslash stateseq}_{2,1}) = 1$$

- using MC

$$\text{\textbackslash statevalfun}_{\pi}(\text{\textbackslash stateseq}_{2,1}) = 1 \quad \text{same as TD}$$

$$\text{\textbackslash statevalfun}_{\pi}(\text{\textbackslash stateseq}_{1,1}) = 2 \quad \text{because it uses } \text{\textbackslash statevalfun}_{\pi}(\text{\textbackslash stateseq}_{1,0})$$

To illustrate the backups:

State_0 --a--> State_1 --a--> ... --a--> State_T (terminal)

MC: Wait for episode to finish, then update $V(s)$
TD: Update $V(s)$ at every transition using $V(s_{\{t+1\}})$

Method	Update Basis	Backup Target
TD(0)	After 1 step, immediate reward + next value	$R_{t+1} + \gamma V(S_{t+1})$
TD(k)	After k steps, sum of k rewards + value at $t + k$	$R_{t+1} + \dots + \gamma^{k-1} R_{t+k}$ + $\gamma^k V(S_{t+k})$
MC (TD(∞) / TD($\lambda = 1$))	End of episode, full sample return	G_t (total return to episode end)

We summarize the comparison of TD and MC:

Feature	Monte Carlo	Temporal Difference
Reliance	Full episode return	Single step + estimated value
Bootstrapping	No	Yes
Update Timing	End of episode	At every step

Bias and Variance

We now have multiple methods for RL

- how can we compare them ?

We will introduce the concepts of Bias and Variance as one metric with which to compare methods.

To be concrete:

What are the advantages/disadvantages of Temporal Difference vs Monte Carlo in updating the Value function ?

In both methods

- the reward that the update to $\text{\statevalfun}_\pi(\text{\stateseq}_{,k})$ depends on *directly*
 - is the immediate one-step reward \rew_{+1}
- subsequent rewards are included via the Value $\text{\statevalfun}_\pi(\text{\stateseq}_{+1})$ of the successor state

Consider iteration k :

In TD (and Dynamic Programming)

- the value of the successor state is *not* the true end-of-episode k return
- it is an *estimate* of the state's value based on the *previous* iteration
- updating one estimate using another estimated is called *bootstrapping*

In MC

- the value of the successor state is the *true* end-of-episode k return

Because TD depends on an estimate

- we say that TD is *biased*.

Moreover, let's suppose that rewards are stochastic

- hence, each reward is a random variable

TD(0) update

- depends on **exactly one** random variable

$$\backslash \text{rew}_{+1}$$

TD(p) update

- depends on $p + 1$ random variables: returns for times +1, …, +p + 1

MC update

- depends on **at least one**

- rewards of *all* subsequent states of the episode π

$$\backslash \text{rew}_{+1}, \dots,$$

Thus, the *variance* of δ

- will be smallest for TD(0)
- will increase for TD(p) as $p > 0$ increases
- will be larger for MC
 - remaining length of episode

But

- the *bias* decreases in the opposite direction of the increase of the variance
- estimate based on more information

Bias and Variance of the methods presented

Method	Bias	Variance	Notes
DP (Dynamic Programming)	Low	Low	Uses full environment model; updates use true expectations.
MC (Monte Carlo)	None	High	Unbiased (targets equal expected return), but episodes can have widely varying returns.
TD (Temporal Difference)	Moderate	Moderate/Low	Bootstrap introduces bias (approximate next value), but reduces variance compared to MC.
V-learning (State Value)	Moderate	Low	Like TD if bootstrapped; low-variance due to value averaging.
Q-learning	Moderate/High	Moderate/High	Off-policy, can have bias (maximization, bootstrapping); variance may increase due to noise in max operator.

Key Points

DP (Dynamic Programming): Lowest bias and variance, but needs a full model (rarely available in practice).

MC: Unbiased estimates, but high variance due to sampling full returns.

TD: Bootstrapping introduces bias but greatly lowers variance versus MC.

V-learning: Similar profile to TD (if bootstrapped); often lower variance than Q-learning due to value averaging.

Q-learning: Can be more biased/variable due to off-policy backups and maximization over sampled estimates.

On-policy versus Off-policy

For many methods (including Q-learning above) the method uses two sub-policies

- the *behavior policy*: the one that chooses an action. For Q-learning this is

With probability ε :

Choose random action a (exploration)

Else:

Choose $a = \text{argmax}_a Q(s, a; \theta)$ (exploitation)

- the *target policy*: the one that we are trying to learn

- target value

$$y_j = r_j + \gamma * \max_{\{a'\}} Q(s'_j, a'; \theta^-) \quad \# \text{ Target Q-value}$$

- the value to which Gradient Descent will guide the NN

- in minimizing Loss

- compares target y_j with prediction $Q_{\text{pred}} = Q(s_j, a_j; \theta)$

When the Behavior and Target policies are identical

- we call the method *On-policy*

If the Behavior and Target polices are potentially different

- we call the method *Off-policy*

Q-learning is an **off-policy** method for several reasons.

The primary reasons are

- target policy is always greedy
- behavior policy
 - ϵ -greedy
 - example from replay buffer may have been conducted with *an older behavior policy* (different NN weights)
 - even if the example's ϵ -greedy choice was from the "greedy" side of the choice

Reason	Behavior Policy (used to collect experience)	Target Policy (used in Q-value update)
Greedy backup (max operator in update)	Behavior policy may choose non-greedy actions due to exploration	Update always targets the action with highest Q-value (optimal policy estimate)
Experience replay buffer	Policy used during sampling in past episodes (possibly older ϵ -greedy or random)	Greedy policy from current network: $a^* = \text{argmax}_a Q(s, a; \theta^-)$
Policy/parameter mismatch over time	Behavior policy determined by weights θ at time of sampling	Target policy determined by weights θ^- at time of update (can differ)

SARSA: On-Policy Q-learning

We present another member of the TD family.

There is a method called *Deep SARSA* that modified Q-learning to be on-policy.

- does not uses Experience Replay buffer
- changes the target from

$$y_j = r_j + \gamma * \max_{\{a'\}} Q(s'_j, a'; \theta^-)$$

- to

$$y_j = r_j + \gamma * Q(s'_j, a'_j; \theta^-)$$

- target uses the action choice a'_j of current behavior policy

$$Q(s'_j, a'_j; \theta^-)$$

- rather than the action that **maximizes** value

$$\max_{\{a'\}} Q(s'_j, a'; \theta^-)$$

Deep SARSA is more stable than DQN

- it is "risk-aware"
 - target is same as behavior when the "exploratory" choice is made for policy
 - so SARSA can learn to avoid risky choices
 - exploratory choices with extreme (negative) rewards
- it is more conservative
 - DQN always makes the greedy choice

Compared to DQN, these characteristics make it

- lower variance
- more likely to converge

But SARSA is higher bias compared to DQN

- the exploratory choice is biased away from the true (optimal) policy

Grid-world: comparing conservative SARSA to risk-loving DQN

We need to navigate in a (4×12) grid

- from start cell S
- to goal cell G
- without "falling off the cliff" (large negative reward: -100) by navigating to cliff cells C

There is a negative reward (-1) for each time step so the reward is maximized by getting to the goal quickly.

.
.
.
S C C C C C C C C G

DQN will favor a path

- that hugs the edge of the cliff
 - faster route to goal G
- but that will fall off the cliff during exploration

SARSA will avoid the cliff

- slower route

```
In [2]: from IPython.display import Image  
Image(filename='images/cliffwalking_paths.gif')
```

```
Out[2]: <IPython.core.display.Image object>
```

Each time a SARSA episode results in falling off a cliff

- a large negative value for $Q(s', a')$ is learned when a' leads to falling off the cliff from state s'
- the large negative reward propagates back through all the preceding (safe) states and actions along the path
 - any path that *could* lead to s' with some probability of a' being chosen *becomes deprecated*
- so policy becomes biased toward moving far from the Cliff
- *even if* the action a' was an *exploratory* action (that won't be repeated)

DQN always assumes the continuation of an episode from the current state is optimal

- the max prevents a' from being chosen in state s'
- so any path that leads to s' *does not become deprecated*

SARSA vs DQN: Summary

Feature	SARSA (On-Policy)	DQN (Off-Policy)
Policy Updated	Behavior policy (ϵ -greedy)	Greedy (optimal) policy
Update Target	$r + \gamma Q(s', a')$ (actual next action)	$r + \gamma \max a' Q(s', a')$ (max over all actions)
Bias/Variance	Higher bias, lower variance	Lower bias, higher variance (risk of instability/overestimation)
Risk Awareness	Accounts for exploration risk (safer in risky environments)	Does not account for risk of exploratory actions
Performance in Large Spaces	Limited without function approx.	Handles large/continuous spaces with neural networks
Experience Replay	Possible but less common	Standard in DQN
Stability	Often more stable, especially in risky/confusing environments	Can be less stable, especially with bad hyperparameters

Learning via Code

Here ([RL_OnPolicy_vs_OffPolicy_code_examples.ipynb](#)) is a notebook that demonstrates

- Q-Learning
 - with and without replay memor
- SARSA

How Each Example Highlights a Salient Point

1. Q-learning (Off-policy):

- **Purpose:** Demonstrates how the behavior policy (epsilon-greedy; explores randomly) differs from the target policy (greedy; always chooses highest Q-value).
- **Salient Point:** Shows that updates use the greedy maximum Q-value (off-policy), even when the sampled action was exploratory.

2. SARSA (On-policy):

- **Purpose:** Demonstrates on-policy learning where both the behavior and target policies are epsilon-greedy and always match.
- **Salient Point:** The Q update uses the value of the exact next action chosen by the behavior policy, making sampling and updating fully aligned.

3. Q-learning with Experience Replay Buffer:

- **Purpose:** Shows how experience transitions are stored in a replay buffer and reused for updates.
- **Salient Point:** Illustrates how off-policy Q-learning leverages exploratory samples and batch updates from the buffer to improve sample efficiency and stabilize learning.

Sources of Variance in RL

We introduced Variance as a way of comparing TD and MC for value-based methods.

But variance is present in many RL methods and is considered a potential impediment to Learning via RL.

- bigger updates in
 - Value function for Value-based methods
 - Parameters for Policy-based methods
- can cause large changes in Policy
- which can lead to unstable training

We list the types of methods affected by each cause.

There multiple sources of Variance, which we summarize below

- stochastic environment
 - stochastic rewards and state transitions
- stochastic policy
- sparse rewards
 - per-episode vs per-step rewards
 - reward estimates become more noisy
- bootstrapping

Source of Variance	Description	Methods Most Affected
Environmental Stochasticity	Random rewards and transitions cause unpredictable outcomes	Monte Carlo, TD, Policy Gradients
Policy Stochasticity	Probabilistic (non-deterministic) action selection	MC, On-policy, Policy Gradients
Long Trajectory Aggregation	Returns summed over many steps compound the randomness	Monte Carlo, n-step methods
Sparse or Delayed Rewards	Few positive signals lead to noisy estimates	Monte Carlo, Value-based
Bootstrapping Error	Using own predictions as targets introduces bias, but lowers variance	TD, Q-Learning
Credit Assignment Difficulty	Uncertainty in linking actions to future rewards increases variance	Policy Gradients, MC
Sample Size & Exploration	Insufficient samples or aggressive exploration cause wide fluctuations	All methods
Model/Optimization Instability	Neural network or optimizer issues amplify variance	Deep RL, Policy Gradient methods
Non-stationarity	Changing environment or policy during training	All methods, esp. online learning

Key Points:

- Variance comes from randomness in data, policy, training procedure, and modeling choices.
- Some sources can be controlled with design choices (e.g., baselines, bootstrapping, averaging), while others are intrinsic to RL problems.

As we introduce new RL methods

- one motivation if to reduce variance

Some potential ways to reduce variance

- $n > 1$ -step ahead Temporal Difference
- estimating updates in mini-batches
 - as in Gradient Descent

In [3]: `print("Done")`

Done

