

Value-based methods: overview

The simplest model-based approach are *Value-based* methods.

They revolve around the idea of

- assigning a *State Value function $\text{\statevalfun}_\pi(\text{\state})$ to each state
 $\text{\state} \in \text{\States}$
$$\text{\statevalfun}_\pi : \text{\States} \rightarrow \text{\Reals}$$
- $\text{\statevalfun}_\pi(\text{\state})$ is an approximation of
$$\text{\E}_\pi(G | \text{\stateseq} = \text{\state})$$

the expected return achievable from state \state

Given statevalfun_π , the optimal deterministic policy is

$$\begin{aligned} & \pi^*(\text{state}) \\ = & \underset{\text{act}}{\text{argmax}} \quad \text{act} \quad \text{statevalfun}_\pi(\text{state}') \\ & \text{transp}(\text{state}', \text{rew} | \text{state}, \text{act}) \neq 0 \end{aligned}$$

- From state state
- Choose the action act
- that results in next state state'
- with maximal $\text{statevalfun}_\pi(\text{state}')$

Note: the argmax results in a *deterministic* policy.

Using a value-based method is practical only if we can discover the state value function $\text{\textbackslash statevalfun}$, which is initially unknown.

Most methods are iterative in nature and create a sequence of increasingly accurate approximations of $\text{\textbackslash statevalfun}$

$$\text{\textbackslash statevalfun}_{\pi,0} \dots \text{\textbackslash statevalfun}_{\pi,k} \dots$$

In the limit

$$\lim_{k \rightarrow \infty} \text{\textbackslash statevalfun}_{\pi,k} = \text{\textbackslash statevalfun}_{\pi}$$

As we obtain an improved approximation $\text{\textbackslash statevalfun}_{\pi,k+1}$

- we may reflect this improved knowledge by updating the policy

Thus, we periodically improve the policy based on improved approximations of
 $\text{\textbackslash statevalfun}_\pi$

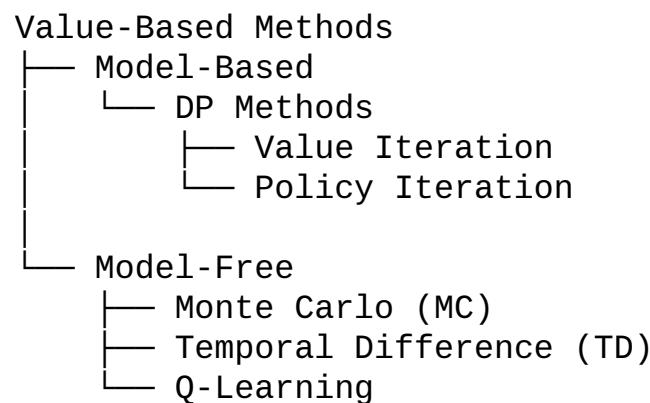
This results in a sequence of increasingly accurate approximations of the policy π

$$\pi_0, \dots, \pi_p, \dots$$

which hopefully converges to π^* .

We will show two broad classes of Value-based methods

- Model-based
- Model-free



A common characteristic of both Model-based and Model-free Value methods

- use of *Dynamic Programming* type solutions.

In Dynamic Programming

- there is a *Bellman update* equation
- which relates the value of a state $\backslash\text{state}$
- to the values of each of the potential successor states $\backslash\text{state}'$
 - which the environment may return in response to the Learner taking an action in state $\backslash\text{state}$

Information flowing from successor to predecessor state is called a *backup*

The Bellman update equation asserts that

$$\text{statevalfun}_\pi(\text{\textbackslash state})$$

can be derived from

- the immediate reward $\text{rew}(\text{\textbackslash state}, \text{\textbackslash act}, \text{\textbackslash state}')$
 - received by taking action $\text{\textbackslash act}$ in state $\text{\textbackslash state}$ and transitioning to state $\text{\textbackslash state}'$
- and the discounted (by γ) value of the successor state $\text{\textbackslash state}'$
 $\text{statevalfun}_\pi(\text{\textbackslash state}')$

The main difference between Model-based and Model-free methods

- a Model-based method does *not need to actively gather* experience
 - if the model is completely known
 - all necessary information is available without actively interacting with the Environment
 - $\text{\textbackslash transp}(\text{\textbackslash state}' | \text{\textbackslash state}, \text{\textbackslash act})$ is known
 $\forall \text{\textbackslash state}, \text{\textbackslash state}' \in \text{\textbackslash States}, \text{\textbackslash act} \in \text{\textbackslash Actions}$
- a Model-free method *must actively interact with the Environment*
 - gathers a *sample* from $\text{\textbackslash transp}(\text{\textbackslash state}' | \text{\textbackslash state}, \text{\textbackslash act})$ for each $\text{\textbackslash state}, \text{\textbackslash act}$ pair in the episode
 - through multiple episodes
 - the *effect* of the true $\text{\textbackslash transp}(\text{\textbackslash state}' | \text{\textbackslash state}, \text{\textbackslash act})$ is derived

Method	Category	Model-based?	Update Equation	Key Characteristics	On-/Off-policy	Convergence Target
Value Iteration	Value-based	Yes	$(V(s) \leftarrow \max_a \sum_{s'} P(s' s, a) [R(s, a, s') + \gamma V(s')])$		Uses full model; synchronous updates	N/A
Policy Iteration	Policy-based	Yes	Eval: $(V^{\pi}(s) \leftarrow \sum_{s'} P(s' s, \pi(s)) [R(s, \pi(s), s') + \gamma V^{\pi}(s')])$ Improve: $(\pi(s) \leftarrow \arg \max_a Q^{\pi}(s, a))$		Alternates evaluation/policy update	N/A

Method	Category	Model-based?	Update Equation	Key Characteristics	On-/Off-policy	Convergence Target
TD(0)	Value-based	No	$(V(s) \leftarrow V(s) + \alpha[r + \gamma V(s') - V(s)])$	Incremental via samples, bootstraps	On-policy	(V^π)
Q-learning	Value-based	No	$(Q(s,a) \leftarrow Q(s,a) + \alpha[r + \gamma \max_{\{a'\}} Q(s',a') - Q(s,a)])$	Off-policy, learns optimal Q	Off-policy	$(Q^{\pi}), (\pi^*)$
Monte Carlo	Value-based	No	$(V(s) \leftarrow V(s) + \alpha[G - V(s)]), \text{ where } (G) = \text{episode return}$	No bootstrapping; full episodes	On-policy	(V^π)

In [2]: `print("Done")`

Done

