

Transformation to add a "missing" numeric feature

Regression: missing feature

We have seen an example of a missing numeric feature in the past.

Recall our example illustrating linear regression

- the first model hypothesized the relationship as

$$\mathbf{y} = \Theta_0 + \Theta_1 \mathbf{x}$$

- Error Analysis revealed a systemic error
- Causing us to add another feature (the square of the first feature)

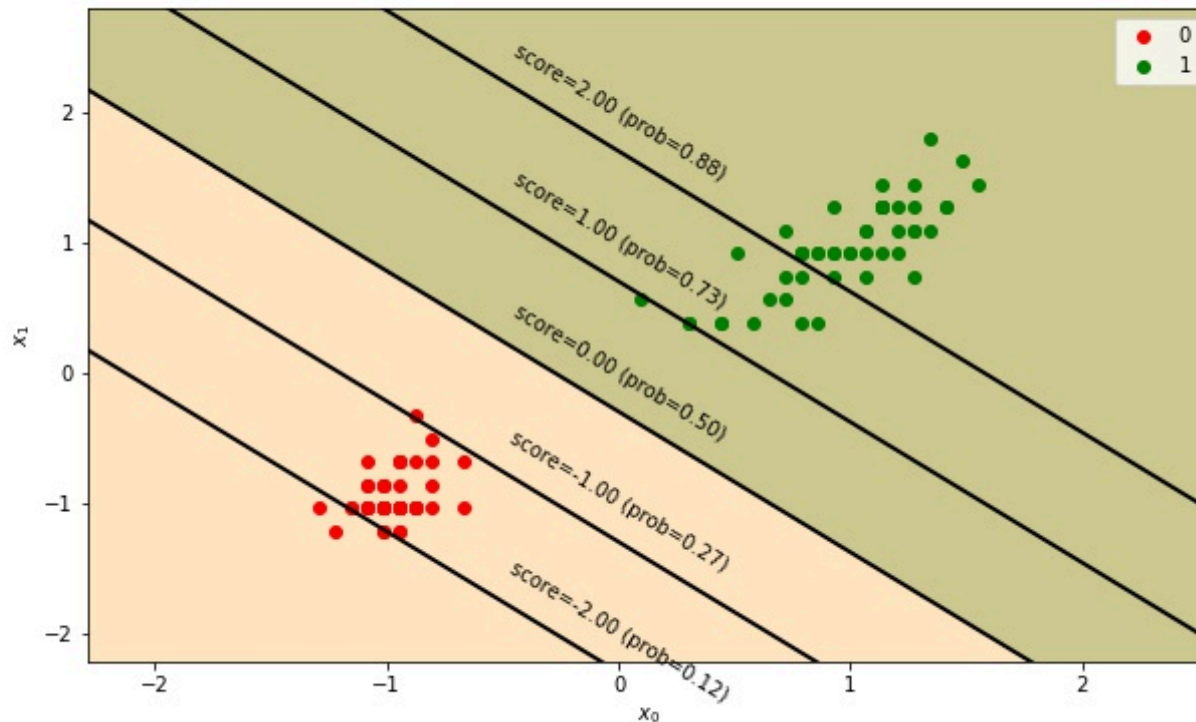
$$\mathbf{y} = \Theta_0 + \Theta_1 \mathbf{x} + \Theta_2 \mathbf{x}^2$$

Classification: missing feature

The Logistic Regression Classifier

- is a type of Classifier
- that creates a *linear surface* to separate classes

Separation bounday as function of probability threshold




But what if the data is such that a linear surface cannot separate classes ?

- we can use a classifier that *does not* assume linear separability (KNN, Decision Trees)
- **or** we can add a feature to make the classes linearly separable
 - here: we illustrate with a numeric feature

Consider Binary Classification on the following "bulls-eye" dataset.


```
In [5]: fig, ax = plt.subplots(1,1, figsize=(6,6) )  
Xc, yc = svmh.make_circles(ax=ax, plot=True)
```



Visually, we can see that the classes are separable, but clearly not by a line.

Here's what one linear classifier (an SVC, which we will study later) produces

```
In [6]: fig, ax = plt.subplots(1,1, figsize=(6,6) )  
svm_clf = svmh.circles_linear(Xc, yc, ax=ax)
```

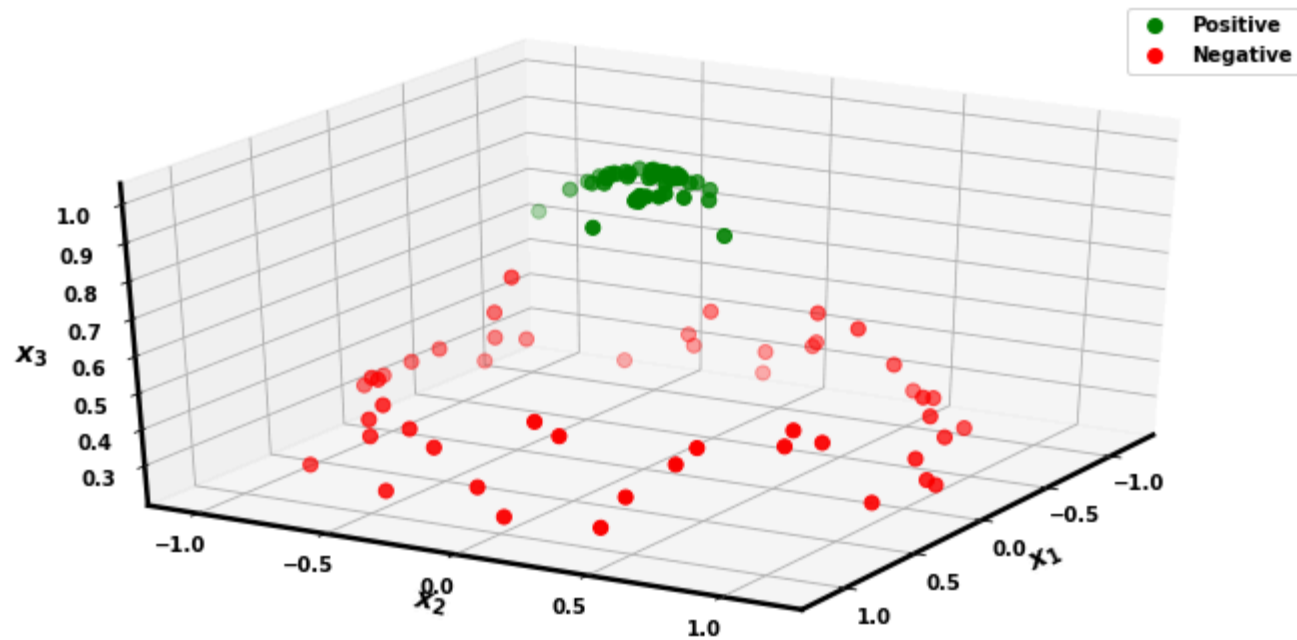


Let's add a new *numeric* feature defined by the (Gaussian) *Radial Basis Function* (RBF)

$$\mathbf{x}_3 = e^{-\sum_j \mathbf{x}_j^2}$$

Our features are now 3 dimensional; let's look at the plot:


```
In [7]: X_w_rbf = svmh.circles_rbf_transform(Xc)
        _ = svmh.plot_3D(X=X_w_rbf, y=yc )
```



Magic !

The new feature is such that it is

- greatest at origin $(\mathbf{x}_1, \mathbf{x}_2) = (0, 0)$
- decreasing as you move away from the origin

The new feature enables a plane that is parallel to the $\mathbf{x}_1, \mathbf{x}_2$ plane to separate the two classes.

We can write the RBF transformation to reference an arbitrary origin \mathbf{x}_c

$$\text{RBF}(\mathbf{x}) = e^{-\|\mathbf{x} - \mathbf{x}_c\|}$$

- $\|\mathbf{x} - \mathbf{x}_c\|$ is a measure of the distance between example \mathbf{x} and reference point \mathbf{x}_c
- In our case
 - $\|\mathbf{x} - \mathbf{x}_c\|$ is the L2 (Euclidean) distance
 - \mathbf{x}_c is the origin $(0, 0)$

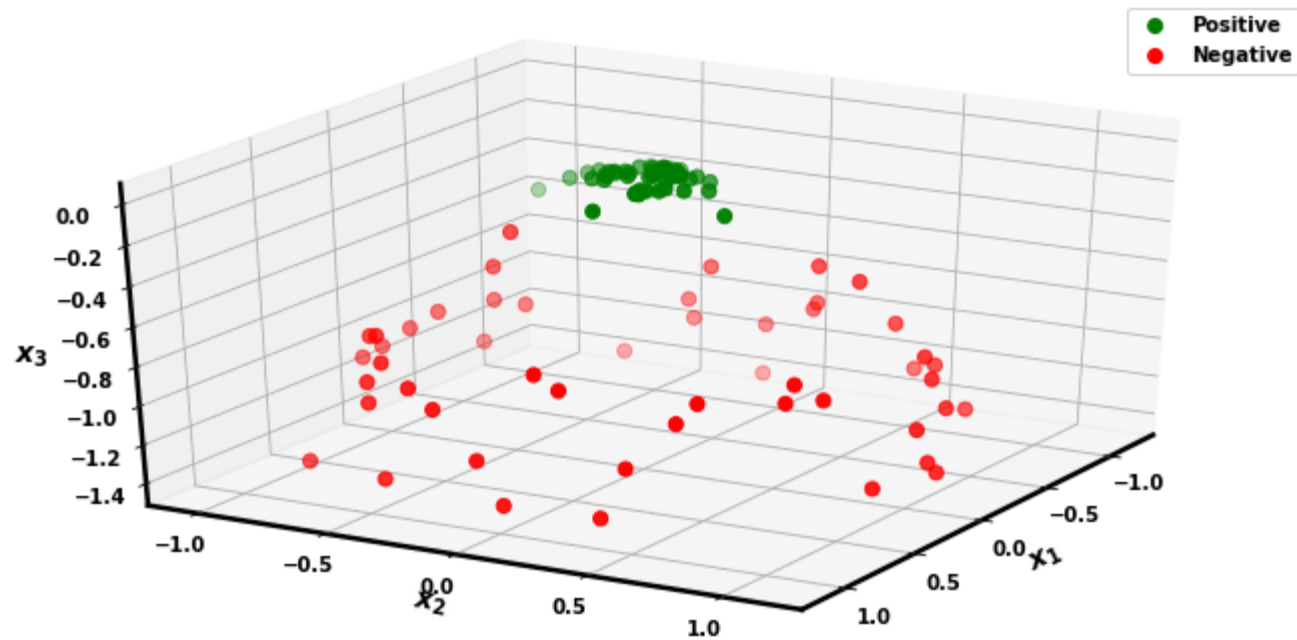
There is an even simpler transformation we could have used

$$\mathbf{x}_3 = -\|\mathbf{x} - \mathbf{x}_c\|^2$$

That is: the (negative) of the L2 distance.

The advantage of the RBF is that it has little effect on points far from the reference point.

```
In [8]: X_w_rad = svmh.circles_radius_transform(Xc)
        _ = svmh.plot_3D(X=X_w_rad, y=yc )
```



The common aspect of each transformation

- observation that there are a set of examples with green labels
- centered around a point (origin)

The transformation added a feature that was greatest in magnitude around those points.mm

Curved boundaries and Linear Classifiers

Recall the transformation of adding a higher order polynomial feature for the "curvy" dataset

$$\mathbf{y} = \Theta_0 + \Theta_1 \mathbf{x} + \Theta_2 \mathbf{x}^2$$

This equation is *still linear* in the two features \mathbf{x}_1 and \mathbf{x}_1^2 .

In Classification, we can create *curved boundaries* that are *still linear* in their features.

- But clearly not linear in raw features

The two plots below use a Classifier requiring Linear Separability of the examples

- the right plot adds a polynomial feature
- creating a curved boundary
- even though the equation is still linear in the features


```

In [9]: svmh = svm_helper.SVM_Helper()

_ = svmh.create_kernel_data()

gamma=1
C=0.1

linear_kernel_svm = svm.SVC(kernel="linear", gamma=gamma)

# Pipelines
feature_map_poly2 = PolynomialFeatures(2)
poly2_approx = pipeline.Pipeline( [ ("feature map", feature_map_poly2),
                                     ("svm", svm.LinearSVC())
                                   ])

classifiers = [ ("SVC", linear_kernel_svm),
                 ("poly (d=2) transform + SVC", poly2_approx)
               ]
_ = svmh.create_kernel_data(classifiers=classifiers)
fig, axs = svmh.plot_kernel_vs_transform()
plt.close()

```


```

/home/kjp/anaconda3/lib/python3.7/site-packages/sklearn/svm/base.py:929: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
  "the number of iterations.", ConvergenceWarning)

```

```
In [10]: fig
```

```
Out[10]:
```



- Left plot shows a boundary that is linear in raw features
- Right plot show a boundary that is linear in transformed features
 - plotted in the dimensions of raw features

The transformation results in a boundary shape with greater flexibility.

Transformations should be motivated by logic, not magic !

Although the transformation on the "bulls-eye" dataset seems magical, we must be skeptical of magic

- There should be some *logical* justification for the added feature
- Without such logic: we are in danger of overfitting and will fail to generalize to test examples

For example:

- Perhaps $\mathbf{x}_1, \mathbf{x}_2$ are geographic coordinates (latitude/longitude)
- There is a distinction (different classes) based on distance from the city center
 $(\mathbf{x}_1, \mathbf{x}_2) = (0, 0)$
 - e.g. Urban/Suburban


Transformation to add a "missing" categorical feature

Recall the dataset where training examples formed two distinct groups

- samples at different points in time

```
In [12]: fig
```

```
Out[12]:
```



How do we pool data that is similar intra-group but different across groups ?

In the above example, it appears that

- The groups are defined by examples gathered at different times: time_0 , time_1
- There is a linear relationship *in each group* in isolation
- There slope of the relationship is *the same* across time
- But the intercept differs across groups
 - Perhaps this reflects a tax or rebate that is independent of price.

If we are correct in hypothesizing that each group is from the same distribution *except for* different intercepts

- the following set of equations describes the data (separately for each of the two groups):

$$\mathbf{y}_{(\text{time}_0)} = \Theta_{(\text{time}_0)} + \Theta_1 * \mathbf{x}$$

$$\mathbf{y}_{(\text{time}_1)} = \Theta_{(\text{time}_1)} + \Theta_1 * \mathbf{x}$$

Trying to fit a line (Linear Regression) as a function of the combined data will be disappointing.

- it will try to force a common intercept

$$\mathbf{y} = \Theta_0 + \Theta_1 * \mathbf{x}$$

when we know that the intercepts are different

We can derive a *single equation* describing both groups by adding

- a Categorical feature Group
- with two possible class values
- indicating which group the example belongs to

Using OHE to encode this categorical feature. we create two binary indicators Is_0, Is_1

$$Is_j^{(i)} = \begin{cases} 1 & \text{if } \mathbf{x}^{(i)} \text{ is in group } j \\ 0 & \text{if } \mathbf{x}^{(i)} \text{ is NOT in group } j \end{cases}$$

To illustrate: for example i in time 0 group, we have

$$Is_0^{(i)} = 1$$

$$Is_1^{(i)} = 0$$

This results in the following equation

$$\mathbf{y} = \Theta_{(\text{time}_0)} * \mathbf{Is}_0 + \Theta_{(\text{time}_1)} * \mathbf{Is}_1 + \Theta_1 * \mathbf{x}$$

Effectively, the equation allows each group to have its own intercept !

- because \mathbf{Is}_0 and \mathbf{Is}_1 are complementary

This transformation caused examples

- that appear different *at the surface level*
- to become *similar* by revealing the *deeper* relationship

Here's what the design matrix \mathbf{X}'' looks like when we add the two indicators:

$$\mathbf{X}'' = \begin{pmatrix} \mathbf{Is}_0 & \mathbf{Is}_1 & \text{other features} \\ 1 & 0 & \dots \\ 0 & 1 & \dots \\ \vdots & & \end{pmatrix} \begin{matrix} \text{time}_0 \\ \text{time}_1 \end{matrix}$$

- Examples from the first time period look similar to the first row
- Examples from the second time period look similar to the second row

Because \mathbf{Is}_0 and \mathbf{Is}_1 are complementary

- we have an instance of the *Dummy Variable Trap*
- we need the usual solution of dropping one binary indicator
 - resulting in

$$\mathbf{y} = \Theta_0 + \Theta'_{(\text{time}_1)} * \mathbf{Is}_1 + \Theta_1 * \mathbf{x}$$

- the intercept term Θ_0 captures the contribution to \mathbf{y} of examples in group 0
- the coefficient $\Theta'_{(\text{time}_1)}$ captures the *incremental* contribution to \mathbf{y} of being in group 1 rather than group 0

Bucketing: making a numeric feature into a Categorical feature

The effect of some features may not have a linear effect on the probability predicted by a Classifier.

Recall the `Age` feature from the Titanic Survival Classification problem.

Our analysis suggested that the Survival probability

- is **not** linear in `Age`

Trying to force `Age` into a linear model would not be appropriate.

Instead, we can create a Categorical synthetic feature AgeBucket

- with class values indicating whether age is in buckets of width 15 years
[0, 15), [15, 30), [30, 45), [45, 60), [60, 75)

Using OHE: we have a binary indicator for each bucket.

This is an example of

- replacing a numeric raw feature
- with a Categorical synthetic feature

to better match the characteristics of the model

Cross features

In our EDA for the Titanic Classification problem we discovered

- being a Female *seemed* to increase the chances of being in the Survived class
- but [deeper analysis \(Classification and Non Numerical Data.ipynb#Conditional-survival-probability-\(condition-on-multiple-attributes\)\)](#) should this to be true *conditional* on not being in Third Class

It seems that we need to identify a group defined by the *intersection* of two conditions

- I_{Female} and $I_{\text{PClass} \neq 3}$

That is, we want to create a feature FNTC (Female Not Third Class)

- that is True
- only for examples whose features are $\text{Sex} = \text{Female}$ and $\text{PClass} \neq 3$

We first create two separate binary features

$\text{Is}_{\text{Female}}$

and

$\text{Is}_{\text{PClass} \neq 3}$

We can create a binary indicator that is the **intersection** of two binary indicators by multiplication

$$I_{\text{FNTC}} = I_{\text{Female}} * I_{\text{PClass} \neq 3}$$

This is called a *cross feature* or a *cross term*.

This cross-feature serves the same purpose as the numeric feature we added to the bulls-eye dataset

- a feature that isolates a subset of examples

In fact, we can use a cross-feature for the bulls-eye dataset

- two binary features
 - one indicating $\mathbf{x}_1^{(i)}$ is close to 0
 - one indicating $\mathbf{x}_2^{(i)}$ is close to 0
- a cross-feature that indicates that $\mathbf{x}^{(i)}$ is close to $(0, 0)$
 - as the product of the two binary features

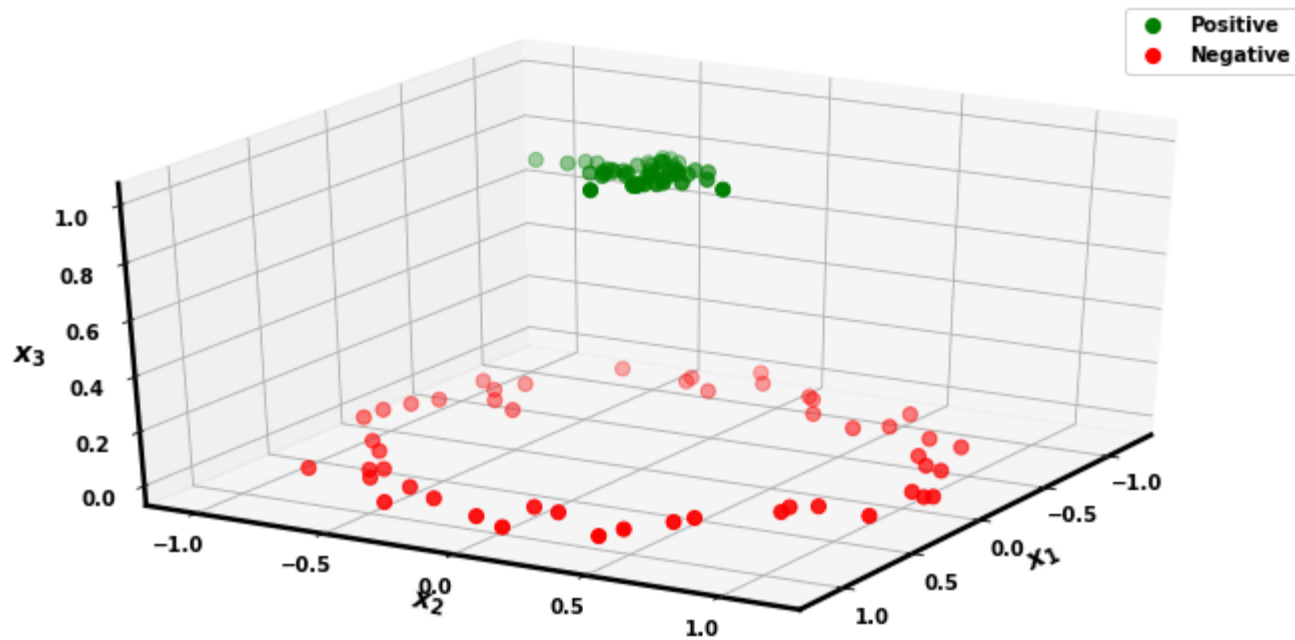
Here we create a cross feature that is `True` if two simpler features hold simultaneously

- $\text{Is}_{\text{near zero } \mathbf{x}_1}$ near zero indicator: $= -\epsilon \leq \mathbf{x}_1 \leq \epsilon$
- $\text{Is}_{\text{near zero } \mathbf{x}_2}$ near zero indicator: $= -\epsilon \leq \mathbf{x}_2 \leq \epsilon$

The cross feature that identifies examples near $(0, 0)$ is

- $\text{Is}_{\text{near}(0,0)} = \text{Is}_{\text{near zero } \mathbf{x}_1} * \text{Is}_{\text{near zero } \mathbf{x}_2}$

```
In [13]: X_w_sq = svmh.circles_square_transform(Xc)
         _ = svmh.plot_3D(X=X_w_sq, y=yc )
```



Cross-features can be abused

Cross terms are very tempting but can be abused when over-used.

- they can be used to identify small subsets of examples for special treatment
- taken to the extreme
 - they can create *one indicator for each training example*
 - essentially: memorizing the training dataset

Memorization of the training set

- usually results in failure to generalize out of sample
- is a hallmark of over-fitting

Here's a picture of the "per example" indicator

First, construct an indicator which is true

- if an example's feature j value is equal to the feature j value of example i :

$$I_{\mathbf{x}_j}^{(i)} = (\mathbf{x}_j = \mathbf{x}_j^{(i)})$$

Now construct a cross feature that combines the indicators for all j and a single example i :

$$I_{\text{example } i} = (\mathbf{x}_1 = \mathbf{x}_1^{(i)}) * (\mathbf{x}_2 = \mathbf{x}_2^{(i)})$$

This cross feature will be true on example i .

We can construct such a cross feature that recognizes any single example.

And here's the design matrix \mathbf{X}'' with a separate intercept per example.

\mathbf{X}'' has m intercept columns, one for each example, forming a diagonal of 1's

$$\mathbf{X}'' = \begin{pmatrix} \mathbf{const} & \text{Is}_{\text{example 1}} & \text{Is}_{\text{example 2}} & \text{Is}_{\text{example 3}} & \dots & \mathbf{other features} \\ 1 & 1 & 0 & 0 & \dots & \\ 1 & 0 & 1 & 0 & \dots & \\ 1 & 0 & 0 & 1 & \dots & \\ \vdots & & & & & \end{pmatrix}$$

We can do the same for $\Theta_1, \Theta_2, \dots, \Theta_n$ resulting in a design matrix \mathbf{X}'' with $m * n$ indicators

- One per example per parameter

$$= \begin{pmatrix} \mathbf{const} & \mathbf{Is}_{\text{example 1}} & (\mathbf{Is}_{\text{example 1}} * \mathbf{x}_1) & (\mathbf{Is}_{\text{example 1}} * \mathbf{x}_2) & \dots & \mathbf{Is}_{\text{example 2}} & (\mathbf{Is}_{\text{example 2}} * \mathbf{x}_1) \\ 1 & 1 & \mathbf{x}_1^{(1)} & \mathbf{x}_2^{(1)} & \dots & 0 & \mathbf{x}_1^{(2)} \\ 1 & 0 & 0 & 0 & \dots & 1 & \mathbf{x}_2^{(2)} \\ \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots \end{pmatrix}$$

Using this as the design matrix in Linear Regression

- Will get a perfect fit to training examples
- Would likely **not generalize** well to out of sample test examples.

When truly justified a small number of complex cross terms are quite powerful.

```
In [14]: print("Done")
```

Done

